

Linear Regression

Data Set

h2o

```
ames <- AmesHousing::make_ames()
ames.h2o <- as.h2o(ames)
```

stratified (*Sale_Price*) training sample

```
set.seed(123)

split <- initial_split(ames, prop = 0.7,
                       strata = "Sale_Price")

ames_train <- training(split)
ames_test <- testing(split)
```

Simple Linear Model

```
model1 <- lm(Sale_Price ~ Gr_Liv_Area, data = ames_train)
```

Fitted regression line (full training)


```
p1 <- model1 %>%
  broom::augment() %>%
  ggplot(aes(Gr_Liv_Area, Sale_Price)) +
  geom_point(size = 1, alpha = 0.3) +
  geom_smooth(se = F, method = "lm") +
  scale_y_continuous(labels = scales::dollar) +
  ggtitle("Fitted regression line")
```

Fitted regression line (restricted range)


```
p2 <- model1 %>%
  broom::augment() %>%
  ggplot(aes(Gr_Liv_Area, Sale_Price)) +
  geom_segment(aes(x = Gr_Liv_Area, y = Sale_Price,
                  xend = Gr_Liv_Area, yend = .fitted),
              alpha = .3) +
  geom_point(size = 1, alpha = 0.3) +
  geom_smooth(se = F, method = "lm") +
  scale_y_continuous(labels = scales::dollar) +
  ggtitle("Fitted regression line (with residuals)")
```

```
grid.arrange(p1, p2, nrow = 2)
```

Fitted regression line



Fitted regression line (with residuals)



```
summary(model1)
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area, data = ames_train)
```

Residuals:

Min	1Q	Median	3Q	Max
-361143	-30668	-2449	22838	331357

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8732.938	3996.613	2.185	0.029 *
Gr_Liv_Area	114.876	2.531	45.385	<0.0000000000000002 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 56700 on 2051 degrees of freedom

Multiple R-squared: 0.5011, Adjusted R-squared: 0.5008

F-statistic: 2060 on 1 and 2051 DF, p-value: < 0.00000000000000022

```
sigma(model1) # RMSE, also Residual Standard Error in summary()
```

```
[1] 56704.78
```

```
sigma(model1)^2 # MSE
```

```
[1] 3215432370
```

Inference

The variability of an estimate is its *standard error (SE)*, the square root of its variance.

t-test for the coefficients are simply the estimated coefficient divided by the standard error (t value = Estimate / Std. Error)

t-test measure the number of standard deviations each coefficient is away from zero (basically $\text{abs}(T) > 2$ is significant at 95% conf)

The confidence interval for coefficients is:

$$\hat{\beta}_j \pm t_{1-\alpha/2, n-p} \hat{SE}(\hat{\beta}_j)$$

```
confint(model1, level = .95)
```

```

                2.5 %      97.5 %
(Intercept) 895.0961 16570.7805
Gr_Liv_Area 109.9121  119.8399
```

Interpretation: We are 95% confident that each one unit increase in Gr_Liv_Area adds between 109.9 and 119.8 dollars to the sale price.

Linear Regression Assumptions:

- 1.) Independent observations
- 2.) The random errors have mean zero, and constant variance
- 3.) The random errors are normally distributed

Multiple Linear Regression

```
(model2 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train))
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train)
```

Coefficients:

```
(Intercept)  Gr_Liv_Area  Year_Built
-2123054.21      99.18      1093.48
```

```
# Equivalent
```

```
(model2 <- update(model1, . ~ . + Year_Built))
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train)
```

Coefficients:

```
(Intercept)  Gr_Liv_Area  Year_Built
-2123054.21      99.18      1093.48
```

```
round(coef(model2), 3)
```

```
(Intercept)  Gr_Liv_Area  Year_Built
-2123054.207      99.176      1093.485
```

```
summary(model3 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built + Gr_Liv_Area:Year_Built, data = ames_train))
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area + Year_Built + Gr_Liv_Area:Year_Built,
    data = ames_train)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-440543  -25191   -1896   17599  281542
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	382194.30149	209192.64043	1.827	0.0678
Gr_Liv_Area	-1483.88103	125.65052	-11.810	<0.0000000000000002
Year_Built	-179.79795	106.40443	-1.690	0.0912
Gr_Liv_Area:Year_Built	0.80371	0.06378	12.601	<0.0000000000000002

```
(Intercept)      .
Gr_Liv_Area      ***
Year_Built       .
Gr_Liv_Area:Year_Built ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 44810 on 2049 degrees of freedom

Multiple R-squared: 0.6887, Adjusted R-squared: 0.6883

F-statistic: 1511 on 3 and 2049 DF, p-value: < 0.00000000000000022

```
round(coef(model3), 3)
```

	(Intercept)	Gr_Liv_Area	Year_Built
	382194.301	-1483.881	-179.798
Gr_Liv_Area:Year_Built	0.804		

```
fit1 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train)
fit2 <- lm(Sale_Price ~ Gr_Liv_Area * Year_Built, data = ames_train)
```

```
# Regression plane
```

```
plot_grid <- expand.grid(
  Gr_Liv_Area = seq(from = min(ames_train$Gr_Liv_Area), to = max(ames_train$Gr_Liv_Area),
    length = 100),
  Year_Built = seq(from = min(ames_train$Year_Built), to = max(ames_train$Year_Built),
    length = 100)
)
```

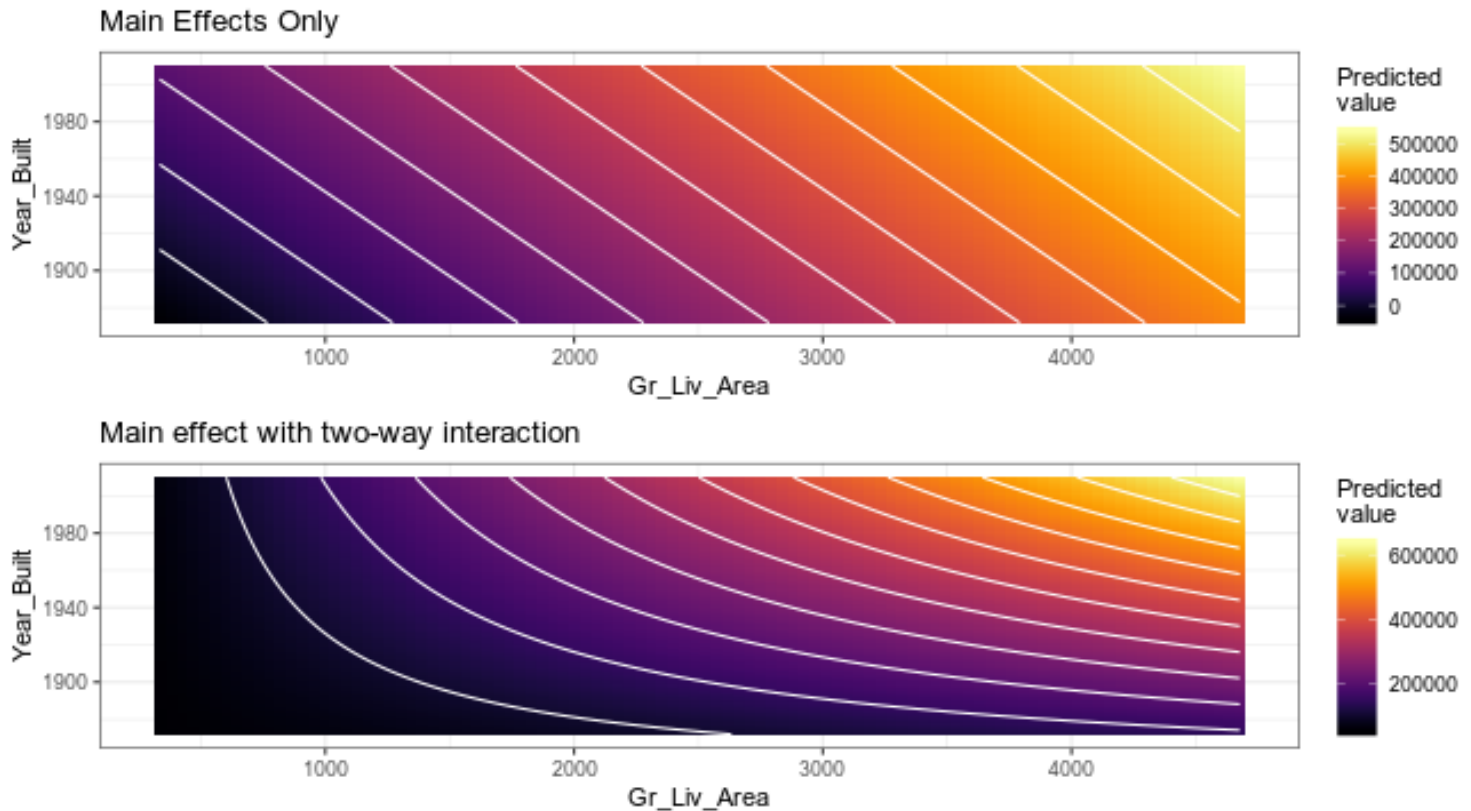
```
plot_grid$y1 <- predict(fit1, newdata = plot_grid)
plot_grid$y2 <- predict(fit2, newdata = plot_grid)
```

```
# Level plots
```

```
p1 <- ggplot(plot_grid, aes(x = Gr_Liv_Area, y = Year_Built,
  z = y1, fill = y1)) +
  geom_tile() +
  geom_contour(color = "white") +
  viridis::scale_fill_viridis(name = "Predicted\nvalue", option = "inferno") +
  theme_bw() +
  ggtitle("Main Effects Only")
```

```
p2 <- ggplot(plot_grid, aes(x = Gr_Liv_Area, y = Year_Built,
  z = y2, fill = y2)) +
  geom_tile() +
  geom_contour(color = "white") +
  viridis::scale_fill_viridis(name = "Predicted\nvalue", option = "inferno") +
  theme_bw() +
  ggtitle("Main effect with two-way interaction")
```

```
gridExtra::grid.arrange(p1, p2, nrow = 2)
```



Full Model

```
model3 <- lm(Sale_Price ~ ., data = ames_train)
broom::tidy(model3)
```

A tibble: 283 x 5

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	-5.61e6	11261881.	-0.498	0.618
2 MS_SubClassOne_Story_1945_and_Older	3.56e3	3843.	0.926	0.355
3 MS_SubClassOne_Story_with_Finished_Atti...	1.28e4	12834.	0.997	0.319
4 MS_SubClassOne_and_Half_Story_Unfinishe...	8.73e3	12871.	0.678	0.498
5 MS_SubClassOne_and_Half_Story_Finished_...	4.11e3	6226.	0.660	0.509
6 MS_SubClassTwo_Story_1946_and_Newer	-1.09e3	5790.	-0.189	0.850
7 MS_SubClassTwo_Story_1945_and_Older	7.14e3	6349.	1.12	0.261
8 MS_SubClassTwo_and_Half_Story_All_Ages	-1.39e4	11003.	-1.27	0.206
9 MS_SubClassSplit_or_Multilevel	-1.15e4	10512.	-1.09	0.276
10 MS_SubClassSplit_Foyer	-4.39e3	8057.	-0.545	0.586

... with 273 more rows

Assessing Model Accuracy

Models 1/2/3:

```
# Train model using 10-fold cross-validation

set.seed(123) # for reproducibility

(cv_model <- train(
  form = Sale_Price ~ Gr_Liv_Area,
  data = ames_train,
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
))
```

Linear Regression

2053 samples
1 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 1846, 1848, 1848, 1848, 1848, 1848, ...

Resampling results:

RMSE	Rsquared	MAE
56410.89	0.5069425	39169.09

Tuning parameter 'intercept' was held constant at a value of TRUE

```
set.seed(123)

cv_model2 <- train(
  Sale_Price ~ Gr_Liv_Area + Year_Built,
  data = ames_train,
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
)

set.seed(123)

suppressWarnings({
  cv_model3 <- train(
    Sale_Price ~ .,
    data = ames_train,
    method = "lm",
    trControl = trainControl(method = "cv", number = 10)
  )
})
```

Accuracy:

```
summary(resamples(list(
  model1 = cv_model,
  model2 = cv_model2,
  model3 = cv_model3
)))
```

Call:

```
summary.resamples(object = resamples(list(model1 = cv_model, model2
  = cv_model2, model3 = cv_model3)))
```

Models: model1, model2, model3

Number of resamples: 10

MAE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
model1	34457.58	36323.74	38943.81	39169.09	41660.81	45005.17	0
model2	28094.79	30594.47	31959.30	32246.86	34210.70	37441.82	0
model3	12458.27	15420.10	16484.77	16258.84	17262.39	19029.29	0

RMSE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
model1	47211.34	52363.41	54948.96	56410.89	60672.31	67679.05	0
model2	37698.17	42607.11	45407.14	46292.38	49668.59	54692.06	0
model3	20844.33	22581.04	24947.45	26098.00	27695.65	39521.49	0

Rsquared

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
model1	0.3598237	0.4550791	0.5289068	0.5069425	0.5619841	0.5965793	0
model2	0.5714665	0.6392504	0.6800818	0.6703298	0.7067458	0.7348562	0
model3	0.7869022	0.9018567	0.9104351	0.8949642	0.9166564	0.9303504	0

Model Concerns / Assumptions**1.) Linear Relationships**

Possible solution to non-linear relationship is by variable transformations:

```
p1 <- ggplot(ames_train, aes(Year_Built, Sale_Price)) +
  geom_point(size = 1, alpha = .4) +
  geom_smooth(se = F) +
  scale_y_continuous(labels = scales::dollar) +
  xlab("Year Built") +
```



```

ggtitle(paste("Non-transformed variables with a \n", "non-linear relationship"))

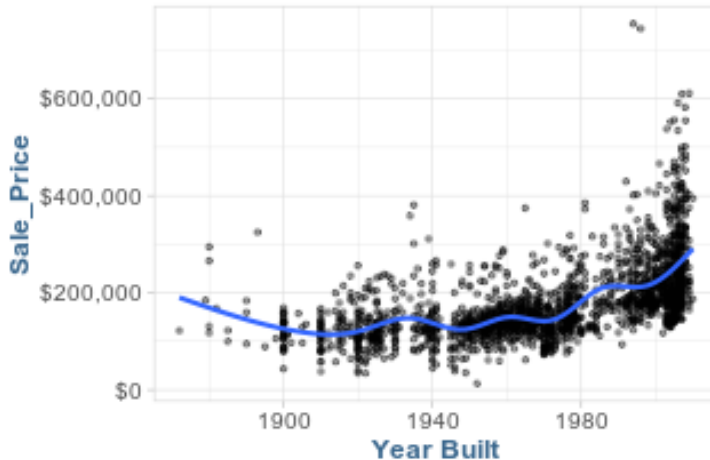
p2 <- ggplot(ames_train, aes(Year_Built, Sale_Price)) +
  geom_point(size = 1, alpha = .4) +
  geom_smooth(se = F, method = "lm") +
  scale_y_log10("Sale Price", labels = scales::dollar, breaks = seq(0, 400000, by = 100000)) +
  xlab("Year Built") +
  ggtitle(paste("Transforming variables can provide a \n", "near-linear relationship"))

gridExtra::grid.arrange(p1, p2, nrow = 1)


`geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'

```

Non-transformed variables with a non-linear relationship



Transforming variables can provide near-linear relationship



2.) Constant variance among residuals

Linear models assume constant variance among error terms (*homoscedasticity*).

Notice the cone shape in Model 1:

```

df1 <- broom::augment(cv_model$finalModel, data = ames_train)


p1 <- ggplot(df1, aes(.fitted, .resid)) +
  geom_point(size = 1, alpha = .4) +
  xlab("Predicted Values") +
  ylab("Residuals") +
  ggtitle("Model 1", subtitle = "Sale_Price ~ Gr_Liv_Area")

df2 <- broom::augment(cv_model3$finalModel, data = ames_train)

```

```
p2 <- ggplot(df2, aes(.fitted, .resid)) +
  geom_point(size = 1, alpha = .4) +
  xlab("Predicted Values") +
  ylab("Residuals") +
  ggtitle("Model 3", subtitle = "Sale_Price ~ .")

gridExtra::grid.arrange(p1, p2, nrow = 1)
```



3.) No autocorrelation

Residuals should be uncorrelated and i.i.d.

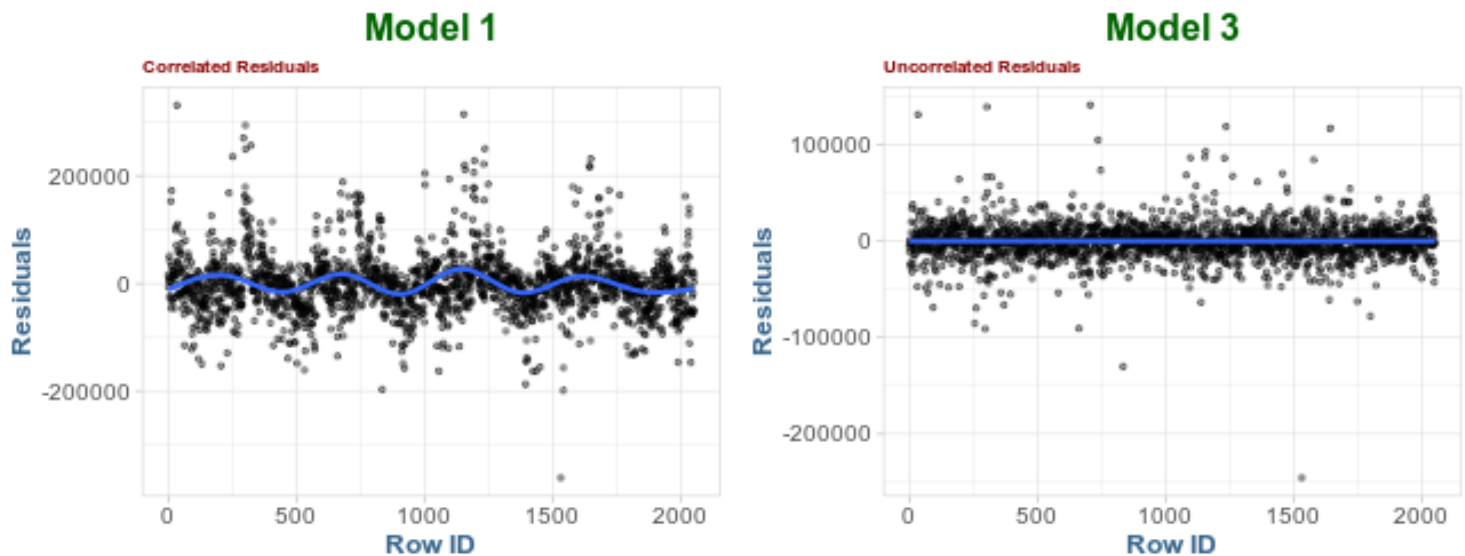
```
df1 <- mutate(df1, id = row_number())
df2 <- mutate(df2, id = row_number())

p1 <- ggplot(df1, aes(id, .resid)) +
  geom_point(size = 1, alpha = .4) +
  geom_smooth(se = F) +
  xlab("Row ID") +
  ylab("Residuals") +
  ggtitle("Model 1", subtitle = "Correlated Residuals")

p2 <- ggplot(df2, aes(id, .resid)) +
  geom_point(size = 1, alpha = .4) +
  geom_smooth(se = F) +
  xlab("Row ID") +
  ylab("Residuals") +
  ggtitle("Model 3", subtitle = "Uncorrelated Residuals")
```

```
gridExtra::grid.arrange(p1, p2, nrow = 1)
```

```
`geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
`geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



Model 1 has a distinct pattern to the residuals (due to being ordered by neighborhood, which is unaccounted for in the model)

4.) More observations than predictors

The number of features cannot exceed the number of observations.

(not a problem in this example)

5.)

No or little multicollinearity.

Collinearity refers to the situation where two or more predictor variables are closely related to one another.

```
cor(ames_train$Garage_Area, ames_train$Garage_Cars)
```

```
[1] 0.8906198
```

Example: *Garage_Area* and *Garage_Cars* are highly correlated.

```
summary(cv_model3) %>%
  broom::tidy() %>%
  filter(term %in% c("Garage_Area", "Garage_Cars"))
```

```
# A tibble: 2 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	Garage_Cars	3021.	1771.	1.71	0.0882
2	Garage_Area	19.7	6.03	3.26	0.00112

Only *Garage_Area* is significant ($p < 0.05$),

However, refit w/o *Garage_Area* term,

```
# model without Garage Area
set.seed(123)
```

```
mod_wo_Garage_Cars <- train(
  Sale_Price ~ .,
  data = select(ames_train, -Garage_Area),
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
)
```

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit may be misleading

Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit

may be misleading

```
summary(mod_wo_Garage_Cars) %>%
  broom::tidy() %>%
  filter(term == "Garage_Cars")
```

```
# A tibble: 1 x 5
  term      estimate std.error statistic    p.value
<chr>      <dbl>     <dbl>     <dbl>    <dbl>
1 Garage_Cars  7161.      1239.       5.78 0.00000000881
```

Garage_Cars becomes significant ($p < 0.001$)

Best to check VIF scores (variance inflation factor)

Principal Component Analysis

PCA -> Linear factorization of features

PCR -> Principal Component Regression

*# performs 10-fold cross validation on a PCR model tuning the
number of principal components to use as predictors from 1-20*


```
set.seed(123)

cv_model_pcr <- train(
  Sale_Price ~ .,
  data = ames_train,
  method = "pcr",
  trControl = trainControl(method = "cv", number = 10),
  preProcess = c("zv", "center", "scale"),
  tuneLength = 20
)

cv_model_pcr$bestTune
```

```
  ncomp
18     18
```

```
ggplot(cv_model_pcr)
```



```
summary(cv_model_pcr)
```

Data: X dimension: 2053 294

Y dimension: 2053 1

Fit method: svdpc

Number of components considered: 18

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps
X	5.987	8.766	11.29	13.39	15.36	17.14	18.58
.outcome	61.142	61.167	62.09	66.28	78.06	78.09	80.87
	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps
X	19.96	21.12	22.23	23.31	24.39	25.41	26.41
.outcome	80.89	81.33	81.39	81.68	81.74	81.89	81.98
	15 comps	16 comps	17 comps	18 comps			
X	27.37	28.30	29.23	30.14			
.outcome	82.01	82.01	82.12	82.28			

Partial least squares

PLS can be viewed as a supervised dimension reduction procedure.

Below is a comparison between PCR and PLS:

```
data(solubility)
```

```
df <- cbind(solTrainX, solTrainY)
```

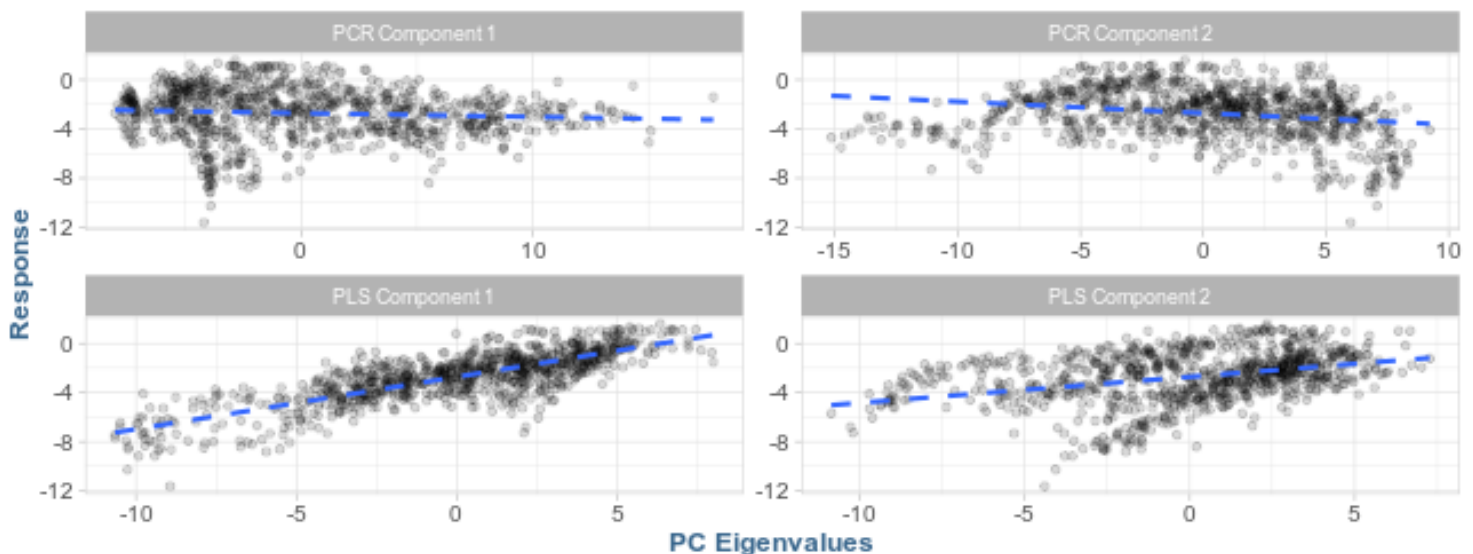
```

pca_df <- recipe(solTrainY ~ ., data = df) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors()) %>%
  step_pca(all_predictors()) %>%
  prep(training = df, retain = T) %>%
  juice() %>%
  select(PC1, PC2, solTrainY) %>%
  rename(`PCR Component 1` = "PC1", `PCR Component 2` = "PC2") %>%
  gather(component, value, -solTrainY)

pls_df <- recipe(solTrainY ~., data = df) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors()) %>%
  step_pls(all_predictors(), outcome = "solTrainY") %>%
  prep(training = df, retain = T) %>%
  juice() %>%
  rename(`PLS Component 1` = "PLS1", `PLS Component 2` = "PLS2") %>%
  gather(component, value, -solTrainY)

pca_df %>%
  bind_rows(pls_df) %>%
  ggplot(aes(value, solTrainY)) +
  geom_point(alpha = .15) +
  geom_smooth(method = "lm", se = F, lty = "dashed") +
  facet_wrap(~ component, scales = "free") +
  labs(x = "PC Eigenvalues", y = "Response")

```



Fit PLS to ames:

```
# perform 10-fold cross validation on a PLS model tuning the
# number of principal components to use as predictors from 1-20
```

```
set.seed(123)
```


```
cv_model_pls <- train(
  Sale_Price ~ .,
  data = ames_train,
  method = "pls",
  trControl = trainControl(method = "cv", number = 10),
  preProcess = c("zv", "center", "scale"),
  tuneLength = 20
)
```

```
# model with lowest RMSE
```

```
cv_model_pls$bestTune
```

```
  ncomp
20     20
```

```
ggplot(cv_model_pls)
```



Feature interpretation

Linear models are monotonic, meaning 1 unit change in X means a constant change in Y.

In PLS we can measure the importance of features by an absolute weighted sum of the regression coefficients.

We can use vip to extract and plot the most important variables:

```
vip(cv_model_pls, num_features = 20, method = "model")
```


Attaching package: 'pls'

The following object is masked from 'package:caret':

R2

The following object is masked from 'package:stats':

loadings



Alternatively, we can use PDP (partial dependence plots) to summarize the relationship.





```
p1 <- pdp::partial(cv_model_pls, pred.var = "Gr_Liv_Area", grid.resolution = 20) %>%
  autoplot() +
  scale_y_continuous(limits = c(0, 300000), labels = scales::dollar)

p2 <- pdp::partial(cv_model_pls, pred.var = "First_Flr_SF", grid.resolution = 20) %>%
  autoplot() +
  scale_y_continuous(limits = c(0, 300000), labels = scales::dollar)

p3 <- pdp::partial(cv_model_pls, pred.var = "Total_Bsmt_SF", grid.resolution = 20) %>%
  autoplot() +
  scale_y_continuous(limits = c(0, 300000), labels = scales::dollar)

p4 <- pdp::partial(cv_model_pls, pred.var = "Garage_Cars", grid.resolution = 4) %>%
  autoplot() +
  scale_y_continuous(limits = c(0, 300000), labels = scales::dollar)
```

```
grid.arrange(p1, p2, p3, p4, nrow = 2)
```



We see all 4 of the top predictive features have a positive linear relationship.

```
# clean up  
rm(list = ls())
```