

Multivariate Adaptive Regression Splines

Data Sets

```
attrition <- attrition %>% mutate_if(is.ordered, factor, order = F)
attrition.h2o <- as.h2o(attrition)

set.seed(123)

ames <- AmesHousing::make_ames()
ames.h2o <- as.h2o(ames)

ames.split <- initial_split(ames, prop = .7, strata = "Sale_Price")

ames.train <- training(ames.split)
ames.test <- testing(ames.split)
```

Overview

Linear models assume the underlying phenomena we are modeling is intrinsically linear which is not usually true. Multivariate adaptive regression splines (MARS) allow us to model non-linear relationships.

Basic strategies for modeling non-linear fits include polynomial regression and step-wise models.

Visually:

```
set.seed(123)

x <- seq(from = 0, to = 2 * pi, length = 500)
y <- sin(x) + rnorm(length(x), sd = .3)

df <- data.table(x, y) %>%
  filter(x < 6)

p1 <- ggplot(df, aes(x, y)) +
  geom_point(alpha = .25) +
  geom_smooth(method = "lm", se = F) +
  ggtitle("(A) Assumed Linear Relationship")

p2 <- ggplot(df, aes(x, y)) +
  geom_point(alpha = .25) +
  geom_smooth(method = "lm", se = F, formula = y ~ poly(x, 2, raw = T)) +
  ggtitle("(B) Degree-2 Polynomial Regression")
```

```

p3 <- ggplot(df, aes(x, y)) +
  geom_point(alpha = .25) +
  geom_smooth(method = "lm", se = F, formula = y ~ poly(x, 3, raw = T)) +
  ggtitle("(C) Degree-3 Polynomial Regression")

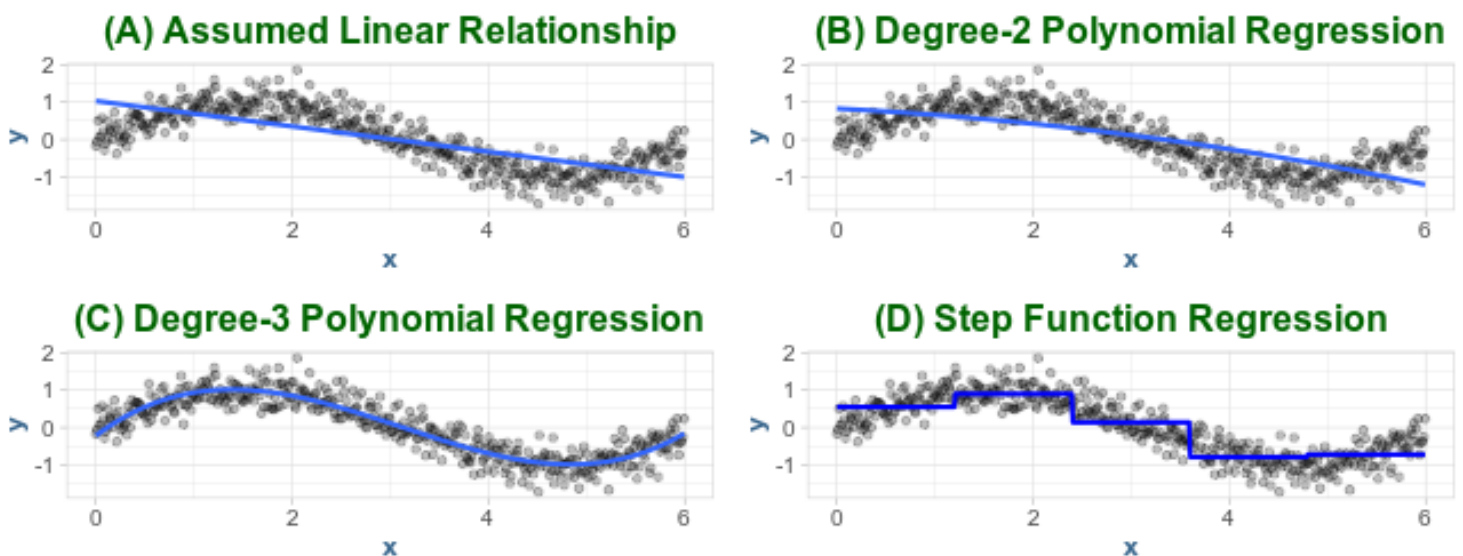
# fit step function model (6 steps)

step_fit <- lm(y ~ cut(x, 5), data = df)
step_pred <- predict(step_fit, df)

p4 <- ggplot(df, aes(x, y)) +
  geom_point(alpha = .25) +
  geom_line(aes(y = step_pred), size = 1, color = "blue") +
  ggtitle("(D) Step Function Regression")

gridExtra::grid.arrange(p1, p2, p3, p4, nrow = 2)

```



Multivariate Adaptive Regression Splines (MARS)

Similar to the step-wise approach, we can access “knots” in the data to model this behavior.

MARS models with *mda* package.

```

mars1 <- mda::mars(df$x, df$y, nk = 3, prune = F)

p1 <- df %>%
  mutate(predicted = as.vector(mars1$fitted.values)) %>%
  ggplot(aes(x, y)) +

```

```
geom_point(size = 1, alpha = .2) +
geom_line(aes(y = predicted), size = 1, color = "blue") +
ggtitle("(A) One Knot")

mars2 <- mda::mars(df$x, df$y, nk = 5, prune = F)

p2 <- df %>%
  mutate(predicted = as.vector(mars2$fitted.values)) %>%
  ggplot(aes(x, y)) +
  geom_point(size = 1, alpha = .2) +
  geom_line(aes(y = predicted), size = 1, color = "blue") +
  ggtitle("(B) Two Knots")

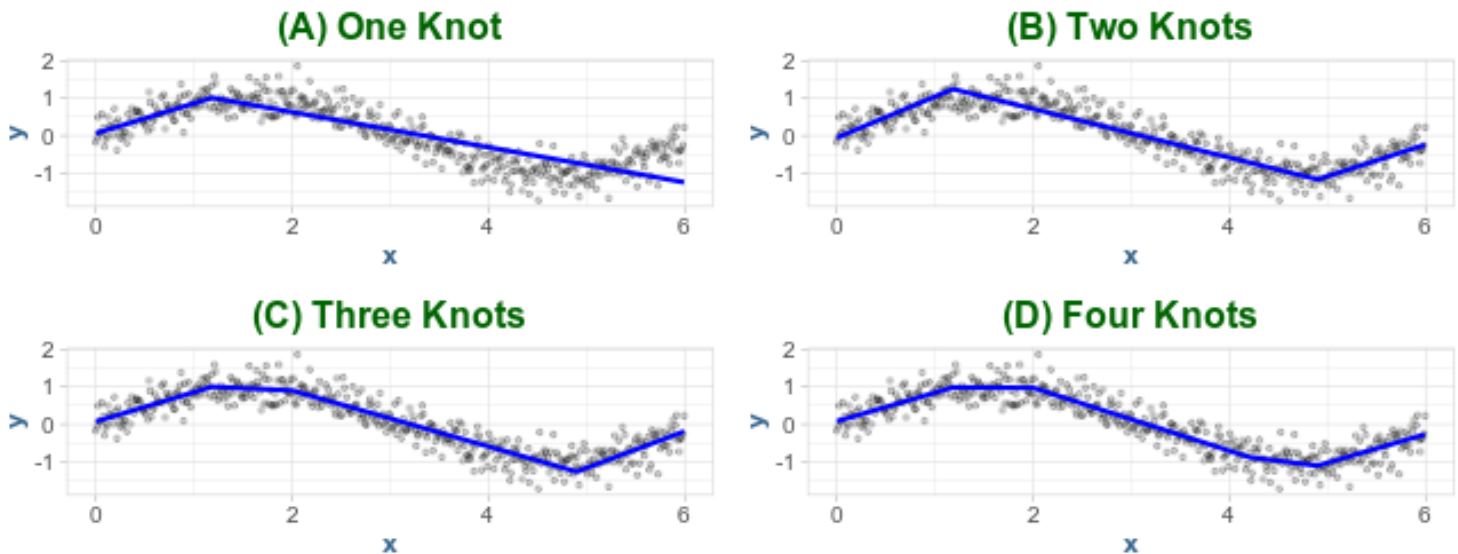
mars3 <- mda::mars(df$x, df$y, nk = 7, prune = F)

p3 <- df %>%
  mutate(predicted = as.vector(mars3$fitted.values)) %>%
  ggplot(aes(x, y)) +
  geom_point(size = 1, alpha = .2) +
  geom_line(aes(y = predicted), size = 1, color = "blue") +
  ggtitle("(C) Three Knots")

mars4 <- mda::mars(df$x, df$y, nk = 9, prune = F)

p4 <- df %>%
  mutate(predicted = as.vector(mars4$fitted.values)) %>%
  ggplot(aes(x, y)) +
  geom_point(size = 1, alpha = .2) +
  geom_line(aes(y = predicted), size = 1, color = "blue") +
  ggtitle("(D) Four Knots")

gridExtra::grid.arrange(p1, p2, p3, p4, nrow = 2)
```



Fitting a basic MARS model

```
mars1 <- earth(
  Sale_Price ~ .,
  data = ames.train
)


print(mars1)
```

```
Selected 36 of 40 terms, and 28 of 307 predictors
Termination condition: RSq changed by less than 0.001 at 40 terms
Importance: Gr_Liv_Area, Year_Built, Total_Bsmt_SF, Overall_QualExcellent, ...
Number of terms at each degree of interaction: 1 35 (additive model)
GCV 547654257    RSS 1047912011488    GRSq 0.9150216    RSq 0.9207205
```

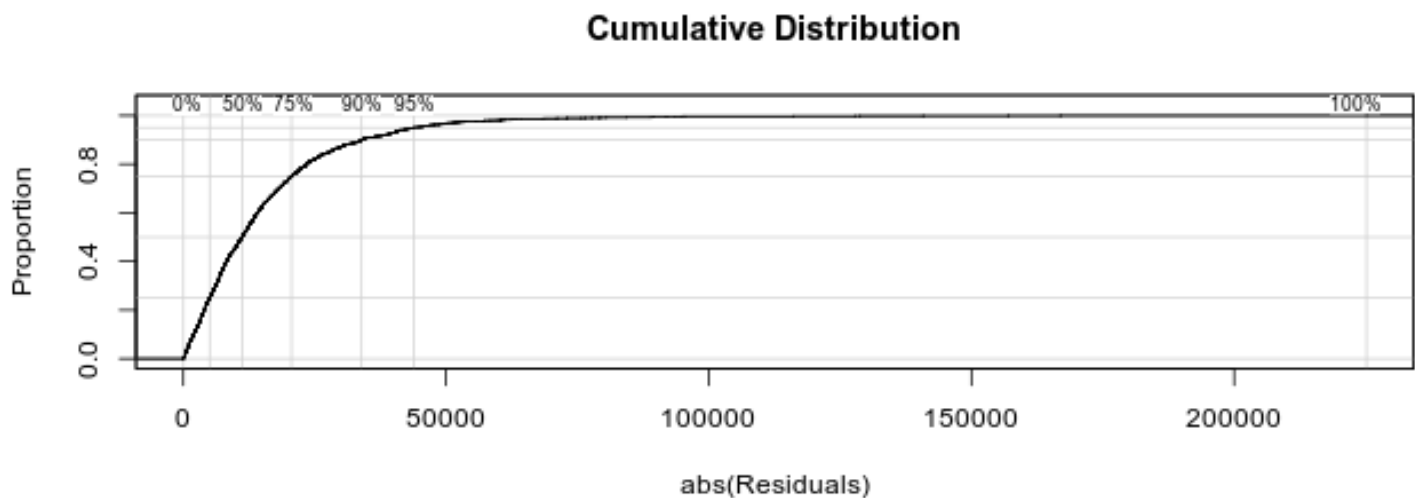
```
summary(mars1) %>% .$coefficients %>% head(10)
```

	Sale_Price
(Intercept)	234803.84945
h(2787-Gr_Liv_Area)	-50.67882
h(Year_Built-2004)	3595.39940
h(2004-Year_Built)	-373.40103
h(Total_Bsmt_SF-1298)	56.30424
h(1298-Total_Bsmt_SF)	-29.95470
h(Bsmt_Unf_SF-536)	-24.47153
h(536-Bsmt_Unf_SF)	16.28784
Overall_QualExcellent	79769.47075
Overall_QualVery_Excellent	117138.64127

```
plot(mars1, which = 1)
```



```
plot(mars1, which = 2)
```



```
mars2 <- earth(
  Sale_Price ~.,
  data = ames.train,
  degree = 2
)

summary(mars2) %>% .$coefficients %>% head(10)
```

```


Sale_Price
(Intercept) 256834.408425081
```

```

h(Gr_Liv_Area-2787)          -144.934131073
h(2787-Gr_Liv_Area)         -49.666036718
h(Year_Built-2004)          4548.490902341
h(2004-Year_Built)          -724.920051940
h(Total_Bsmt_SF-1298)       80.606680445
h(1298-Total_Bsmt_SF)       -41.967572528
h(Bsmt_Unf_SF-1017)*h(2787-Gr_Liv_Area) -0.023524460
h(1017-Bsmt_Unf_SF)*h(2787-Gr_Liv_Area)  0.008535269
Condition_1Norm*h(Gr_Liv_Area-2787)      278.504549507

```

```
plot(mars2, which = 1)
```



Tuning

As always, we will use a cross-validated grid search procedure to tune the hyperparameters.

First pass:

```

hyper.grid <- expand.grid(
  degree = 1:3,
  nprune = seq(2, 100, length.out = 10) %>% floor()
)

```

```
head(hyper.grid)
```

	degree	nprune
1	1	2
2	2	2
3	3	2
4	1	12

```
5      2      12
6      3      12
```

```
set.seed(123)

suppressWarnings(print({
cv.mars <- train(
  x = ames.train %>% select(-Sale_Price),
  y = ames.train$Sale_Price,
  method = "earth",
  metric = "RMSE",
  trControl = trainControl(method = "cv", number = 10),
  tuneGrid = hyper.grid
})))
```

Multivariate Adaptive Regression Spline

2053 samples
80 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 1846, 1848, 1848, 1848, 1848, 1848, ...

Resampling results across tuning parameters:

degree	nprune	RMSE	Rsquared	MAE
1	2	56943.09	0.4964264	39557.80
1	12	31634.16	0.8447279	20207.96
1	23	28407.87	0.8732735	17936.68
1	34	28235.91	0.8753639	17250.10
1	45	28140.74	0.8759693	17065.57
1	56	28140.74	0.8759693	17065.57
1	67	28140.74	0.8759693	17065.57
1	78	28140.74	0.8759693	17065.57
1	89	28140.74	0.8759693	17065.57
1	100	28140.74	0.8759693	17065.57
2	2	56150.99	0.5111693	39532.38
2	12	31572.67	0.8450372	20993.65
2	23	29873.41	0.8629176	17870.00
2	34	28410.99	0.8722307	16630.36
2	45	28051.85	0.8752016	16189.05
2	56	27899.28	0.8762783	16171.49
2	67	27899.28	0.8762783	16171.49
2	78	27899.28	0.8762783	16171.49
2	89	27899.28	0.8762783	16171.49
2	100	27899.28	0.8762783	16171.49

3	2	56571.66	0.5061590	39916.57
3	12	33411.42	0.8279242	22001.04
3	23	30829.04	0.8545617	18542.74
3	34	29795.30	0.8621228	17312.10
3	45	29288.71	0.8667649	16845.19
3	56	29322.13	0.8665347	16811.13
3	67	29322.13	0.8665347	16811.13
3	78	29322.13	0.8665347	16811.13
3	89	29322.13	0.8665347	16811.13
3	100	29322.13	0.8665347	16811.13

RMSE was used to select the optimal model using the smallest value.
The final values used for the model were `nprune = 56` and `degree = 2`.

```
cv.mars$bestTune
```

```
  nprune degree
16      56      2
```


```
cv.mars$results %>%
  as_tibble() %>%
  arrange(RMSE)
```

```
# A tibble: 30 x 8
```

	degree	nprune	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	2	56	27899.	0.876	16171.	15103.	0.129	2354.
2	2	67	27899.	0.876	16171.	15103.	0.129	2354.
3	2	78	27899.	0.876	16171.	15103.	0.129	2354.
4	2	89	27899.	0.876	16171.	15103.	0.129	2354.
5	2	100	27899.	0.876	16171.	15103.	0.129	2354.
6	2	45	28052.	0.875	16189.	15525.	0.132	2378.
7	1	45	28141.	0.876	17066.	8845.	0.0757	1915.
8	1	56	28141.	0.876	17066.	8845.	0.0757	1915.
9	1	67	28141.	0.876	17066.	8845.	0.0757	1915.
10	1	78	28141.	0.876	17066.	8845.	0.0757	1915.

```
# ... with 20 more rows
```

```
ggplot(cv.mars) +
  scale_y_continuous(labels = scales::comma)
```

Refinement:

```
refine.grid <- expand.grid(
  degree = 1,
  nprune = seq(from = 30, to = 45)
)

suppressWarnings(print({
  cv.mars2 <- train(
    x = ames.train %>% select(-Sale_Price),
    y = ames.train$Sale_Price,
    method = "earth",
    metric = "RMSE",
    trControl = trainControl(method = "cv", number = 10),
    tuneGrid = refine.grid
  )
}))
```

Multivariate Adaptive Regression Spline

2053 samples
80 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 1846, 1848, 1848, 1847, 1848, 1849, ...

Resampling results across tuning parameters:

nprune	RMSE	Rsquared	MAE
30	28037.08	0.8768351	17239.50

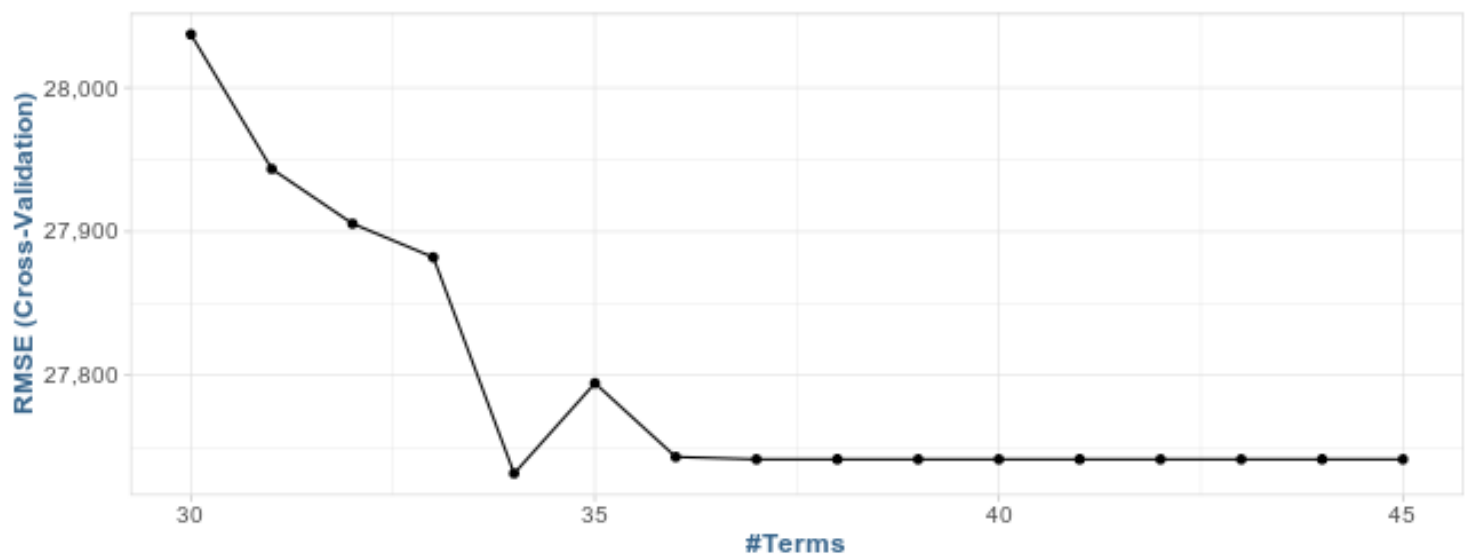
31	27943.54	0.8777220	17170.66
32	27905.51	0.8780454	17059.47
33	27882.12	0.8786504	17016.62
34	27731.86	0.8798819	16922.83
35	27794.61	0.8793995	16910.55
36	27743.34	0.8798244	16872.24
37	27741.71	0.8798384	16872.96
38	27741.71	0.8798384	16872.96
39	27741.71	0.8798384	16872.96
40	27741.71	0.8798384	16872.96
41	27741.71	0.8798384	16872.96
42	27741.71	0.8798384	16872.96
43	27741.71	0.8798384	16872.96
44	27741.71	0.8798384	16872.96
45	27741.71	0.8798384	16872.96

Tuning parameter 'degree' was held constant at a value of 1
 RMSE was used to select the optimal model using the smallest value.
 The final values used for the model were nprune = 34 and degree = 1.

```
cv.mars2$bestTune
```

```
nprune degree
5      34      1
```

```
ggplot(cv.mars2) +
  scale_y_continuous(labels = scales::comma)
```



Feature Interpretation

Variable Importance Plots

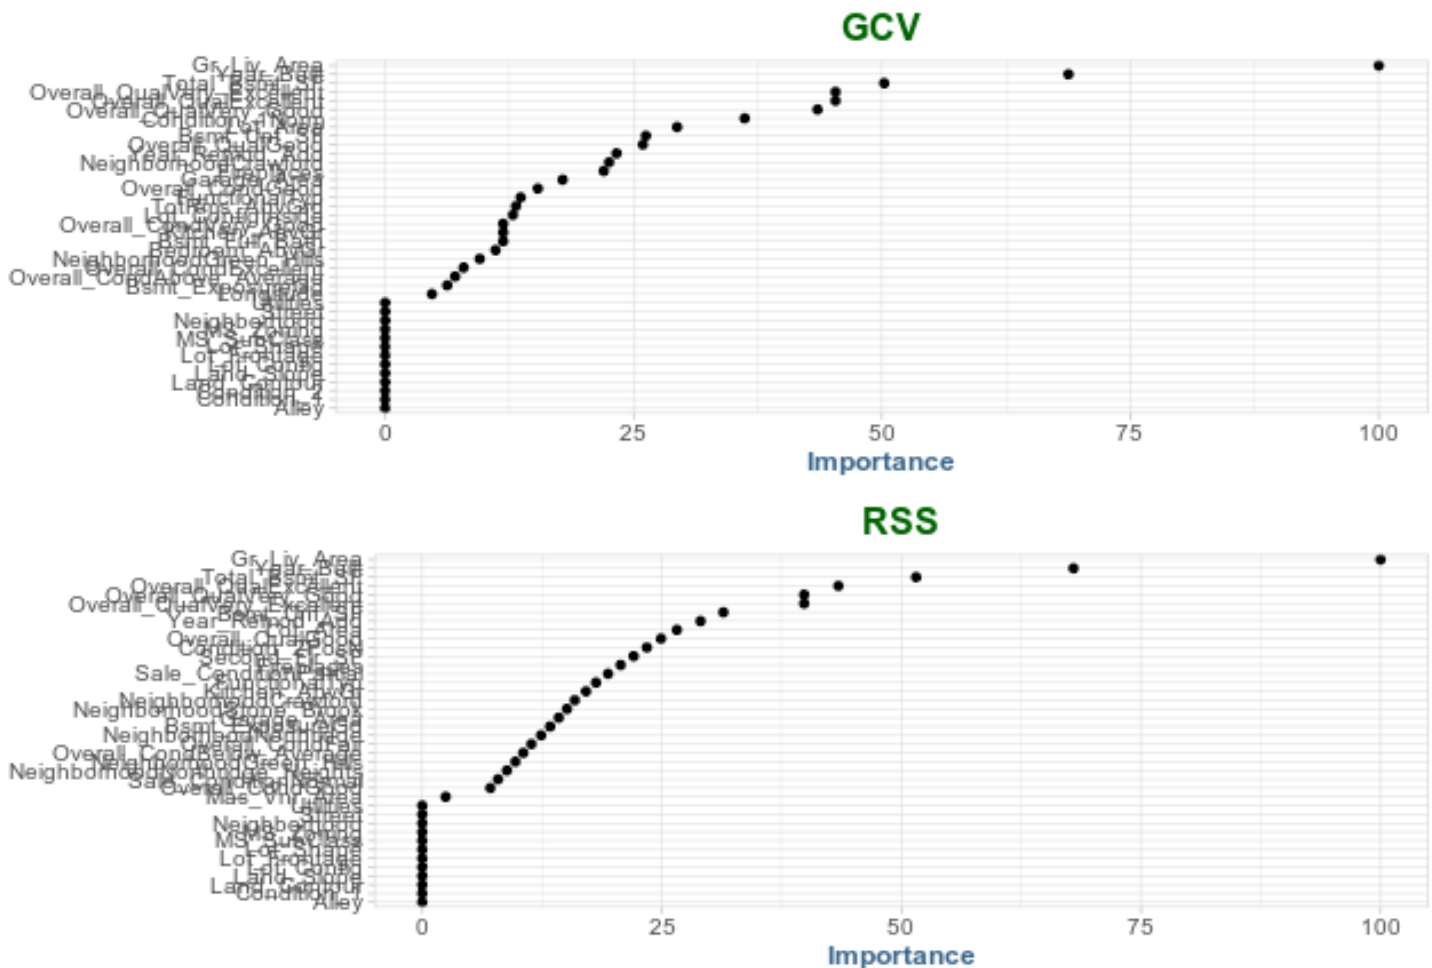
```
p1 <- vip::vip(cv.mars, num_features = 40, bar = F, value = "gcv") +  
  ggtitle("GCV")
```

Warning in vip.default(cv.mars, num_features = 40, bar = F, value = "gcv"): The ``bar`` argument has been deprecated in favor of the new ``geom`` argument. It will be removed in version 0.3.0.

```
p2 <- vip::vip(cv.mars2, num_features = 40, bar = F, value = "rss") +  
  ggtitle("RSS")
```

Warning in vip.default(cv.mars2, num_features = 40, bar = F, value = "rss"): The ``bar`` argument has been deprecated in favor of the new ``geom`` argument. It will be removed in version 0.3.0.

```
gridExtra::grid.arrange(p1, p2, nrow = 2)
```



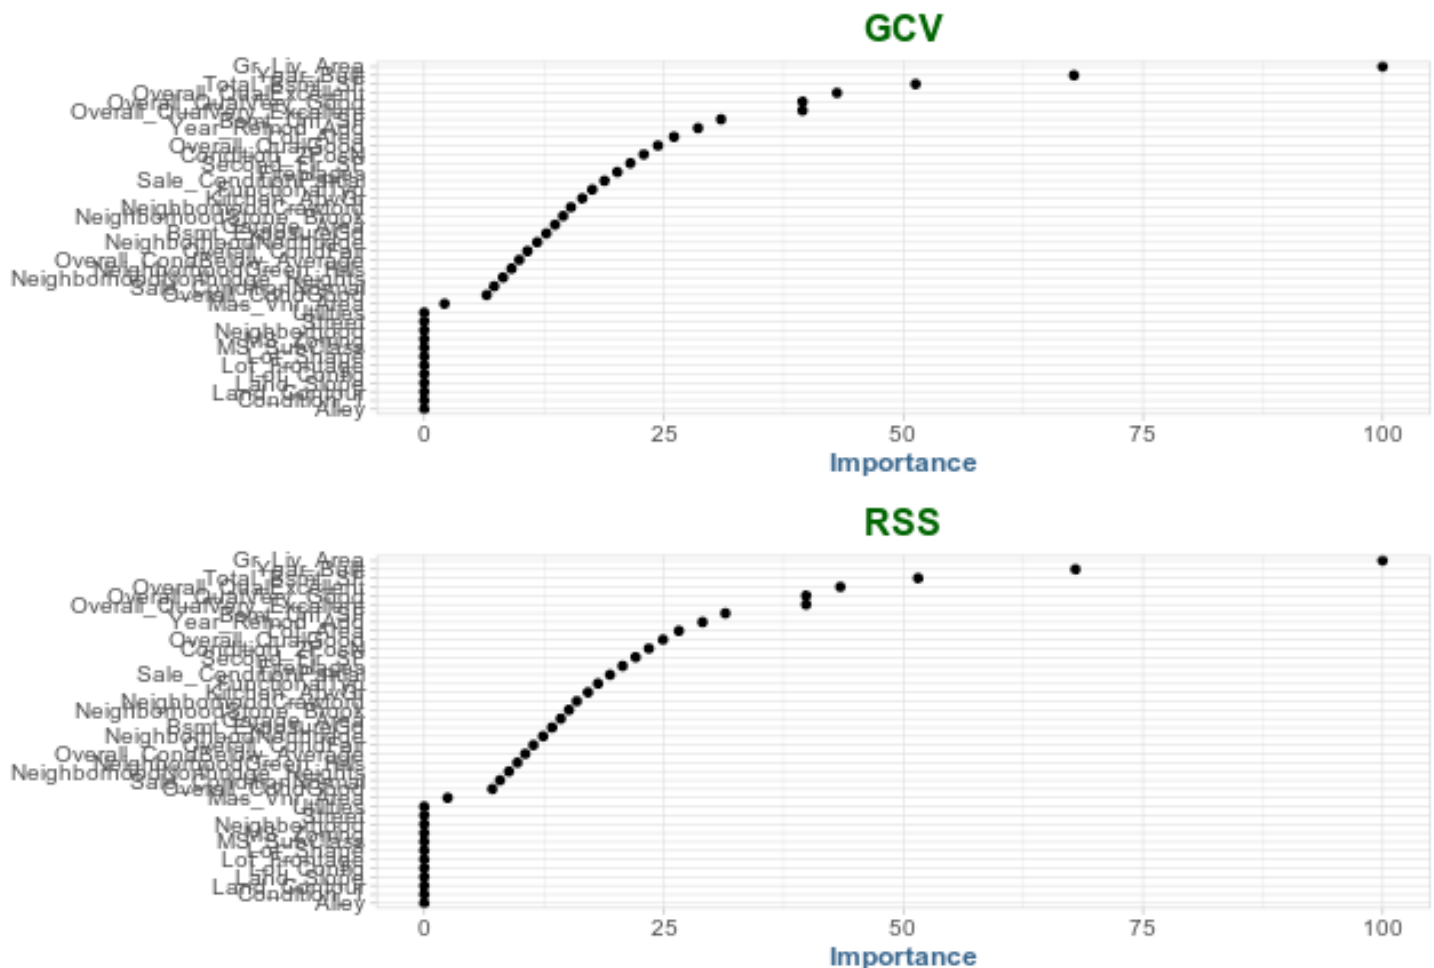
```
p1 <- vip::vip(cv.mars2, num_features = 40, bar = F, value = "gcv") +
  ggtitle("GCV")
```

Warning in vip.default(cv.mars2, num_features = 40, bar = F, value = "gcv"): The `bar` argument has been deprecated in favor of the new `geom` argument. It will be removed in version 0.3.0.

```
p2 <- vip::vip(cv.mars2, num_features = 40, bar = F, value = "rss") +
  ggtitle("RSS")
```

Warning in vip.default(cv.mars2, num_features = 40, bar = F, value = "rss"): The `bar` argument has been deprecated in favor of the new `geom` argument. It will be removed in version 0.3.0.

```
gridExtra::grid.arrange(p1, p2, nrow = 2)
```



```
# extract coefficients, covert to tidy & filter for interaction
```

```
cv.mars2$finalModel %>%
  coef() %>%
```

```
broom::tidy()
```

```
Warning: 'tidy.numeric' is deprecated.
See help("Deprecated")
```

```
# A tibble: 34 x 2
```

	names	x
	<chr>	<dbl>
1	(Intercept)	238598.
2	h(2787-Gr_Liv_Area)	-50.5
3	h(Year_Built-2004)	3735.
4	h(2004-Year_Built)	-336.
5	h(Total_Bsmt_SF-1298)	56.8
6	h(1298-Total_Bsmt_SF)	-28.5
7	h(Bsmt_Unf_SF-536)	-24.2
8	h(536-Bsmt_Unf_SF)	15.4
9	Overall_QualExcellent	80797.
10	Overall_QualVery_Excellent	119191.

```
# ... with 24 more rows
```

```
# Construct partial dependence plots
```

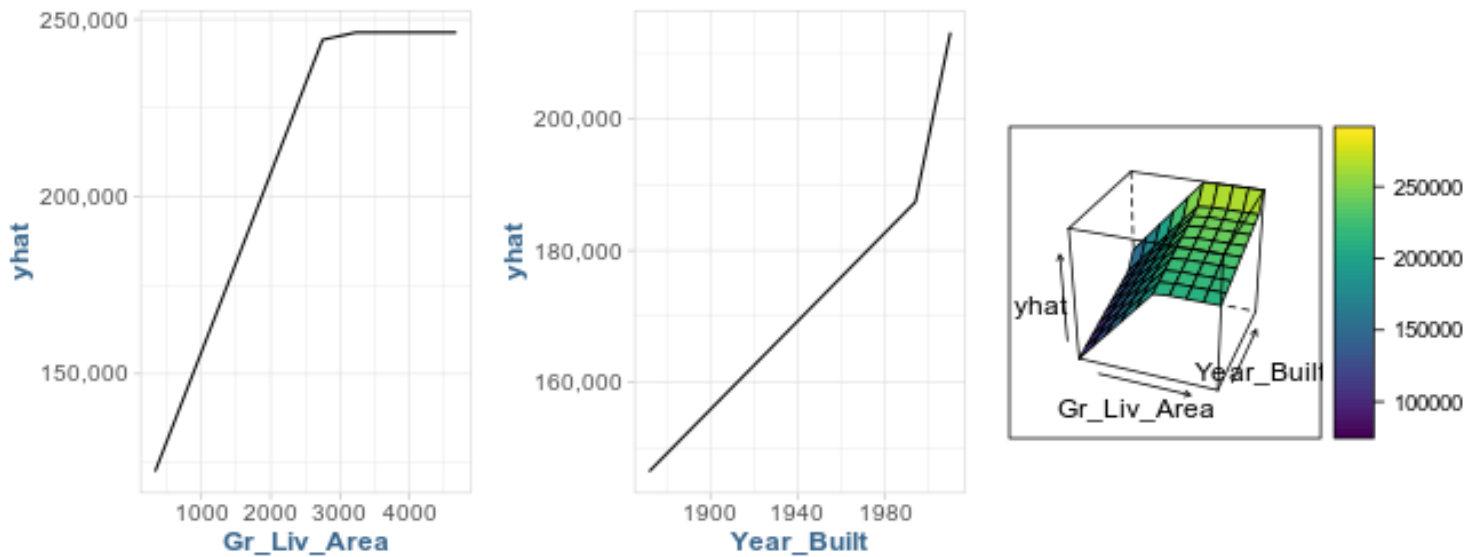
```
p1 <- pdp::partial(cv.mars2, pred.var = "Gr_Liv_Area", grid.resolution = 10) %>%
  autoplot() +
  scale_y_continuous(labels = scales::comma)
```

```
p2 <- pdp::partial(cv.mars2, pred.var = "Year_Built", grid.resolution = 10) %>%
  autoplot() +
  scale_y_continuous(labels = scales::comma)
```

```
p3 <- pdp::partial(cv.mars2, pred.var = c("Gr_Liv_Area", "Year_Built"),
  grid.resolution = 10) %>%
  plotPartial(levelplot = FALSE, zlab = "yhat", drape = TRUE, colorkey = TRUE,
    screen = list(z = -20, x = -60))
```

```
# Display plots side by side
```

```
gridExtra::grid.arrange(p1, p2, p3, ncol = 3)
```



Attrition data

```
df <- rsample::attrition %>% mutate_if(is.ordered, factor, order = F)

# Create training (70%) and test (30%) sets for the attrition data.
set.seed(123)

churn.split <- initial_split(df, prop = .7, strata = "Attrition")
churn.train <- training(churn.split)
churn.test <- testing(churn.split)

set.seed(123)

suppressWarnings(print({
tuned.mars <- train(
  x = subset(churn.train, select = -Attrition),
  y = churn.train$Attrition,
  method = "earth",
  trControl = trainControl(method = "cv", number = 10),
  tuneGrid = hyper.grid
})))
```

Multivariate Adaptive Regression Spline

```
1030 samples
 30 predictor
 2 classes: 'No', 'Yes'
```

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 926, 926, 927, 928, 928, 927, ...

Resampling results across tuning parameters:

degree	nprune	Accuracy	Kappa
1	2	0.8378862	0.005358948
1	12	0.8708500	0.401114326
1	23	0.8661373	0.430216328
1	34	0.8641099	0.428656510
1	45	0.8621777	0.426509019
1	56	0.8621777	0.426509019
1	67	0.8621777	0.426509019
1	78	0.8621777	0.426509019
1	89	0.8621777	0.426509019
1	100	0.8621777	0.426509019
2	2	0.8417225	0.224053300
2	12	0.8446920	0.299911342
2	23	0.8437495	0.332054852
2	34	0.8389424	0.325019482
2	45	0.8408937	0.342780742
2	56	0.8437503	0.359336395
2	67	0.8437503	0.359336395
2	78	0.8437503	0.359336395
2	89	0.8437503	0.359336395
2	100	0.8437503	0.359336395
3	2	0.8514312	0.266403262
3	12	0.8436073	0.280383943
3	23	0.8232658	0.251074784
3	34	0.8116433	0.231032715
3	45	0.8165072	0.237656881
3	56	0.8184490	0.250682779
3	67	0.8184490	0.250682779
3	78	0.8184490	0.250682779
3	89	0.8184490	0.250682779
3	100	0.8184490	0.250682779

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were nprune = 12 and degree = 1.


```
tuned.mars$bestTune
```

```

nprune degree
2      12      1

```

```
ggplot(tuned.mars)
```



```
# train logistic regression model
set.seed(123)
```

```
glm.mod <- train(
  Attrition ~ .,
  data = churn.train,
  method = "glm",
  family = "binomial",
  preProc = c("zv", "center", "scale"),
  trControl = trainControl(method = "cv", number = 10)
)
```

```
# train regularized logistic regression model
set.seed(123)
```

```
penalized.mod <- train(
  Attrition ~ .,
  data = churn.train,
  method = "glmnet",
  family = "binomial",
  preProc = c("zv", "center", "scale"),
  trControl = trainControl(method = "cv", number = 10),
  tuneLength = 10
)
```

```
# extract out of sample performance measures
summary(resamples(list(
```


Table 1: Cross-validated accuracy results for tuned MARS and regression models.

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Logistic_model	0.8365385	0.8495146	0.8792476	0.8757893	0.8907767	0.9313725	0
Elastic_net	0.8446602	0.8759280	0.8834951	0.8835759	0.8915469	0.9411765	0
MARS_model	0.8155340	0.8578463	0.8780697	0.8708500	0.8907767	0.9029126	0

```
Logistic_model = glm.mod,  
Elastic_net = penalized.mod,  
MARS_model = tuned.mars  
)))$statistics$Accuracy %>%  
kableExtra::kable(caption = "Cross-validated accuracy results for tuned MARS and regression models",  
kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))  
  
# clean up  
rm(list = ls())
```