

Chapter 8

Lab

Fitting Classification Trees

```
carseats <- as.data.table(ISLR::Carseats)

carseats[, High := as.factor(ifelse(Sales <= 8, "No", "Yes"))]

tree.carseats <- tree(formula = High ~ .-Sales, data = carseats)

summary(tree.carseats)
```

Classification tree:

```
tree(formula = High ~ . - Sales, data = carseats)
```

Variables actually used in tree construction:

```
[1] "ShelveLoc" "Price" "Income" "CompPrice" "Population"
```

```
[6] "Advertising" "Age" "US"
```

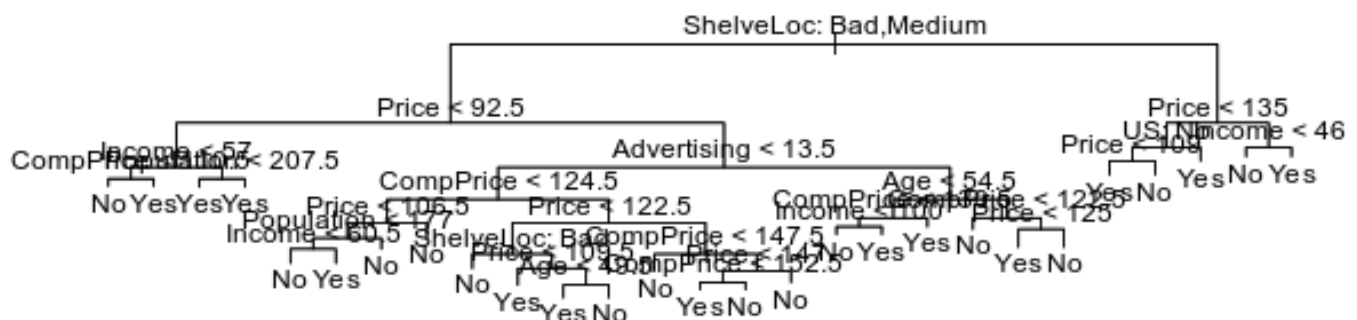
Number of terminal nodes: 27

Residual mean deviance: 0.4575 = 170.7 / 373

Misclassification error rate: 0.09 = 36 / 400

```
plot(tree.carseats)
```

```
text(tree.carseats, pretty = 0)
```



```

set.seed(2)

train <- sample(1:nrow(carseats), 200)

carseats.test <- carseats[-train]
high.test <- carseats[-train]$High

tree.carseats <- tree(High ~.-Sales, data = carseats, subset = train)
tree.pred <- predict(tree.carseats, carseats.test, type = "class")

table(tree.pred, high.test)

```

```

      high.test
tree.pred No Yes
      No  104  33
      Yes   13  50

```

```

set.seed(3)

cv.carseats <- cv.tree(tree.carseats, FUN = prune.misclass)
names(cv.carseats)

```

```
[1] "size"    "dev"     "k"       "method"
```

```
cv.carseats
```

```
$size
```

```
[1] 21 19 14  9  8  5  3  2  1
```

```
$dev
```

```
[1] 74 76 81 81 75 77 78 85 81
```

```
$k
```

```
[1] -Inf  0.0  1.0  1.4  2.0  3.0  4.0  9.0 18.0
```

```
$method
```

```
[1] "misclass"
```

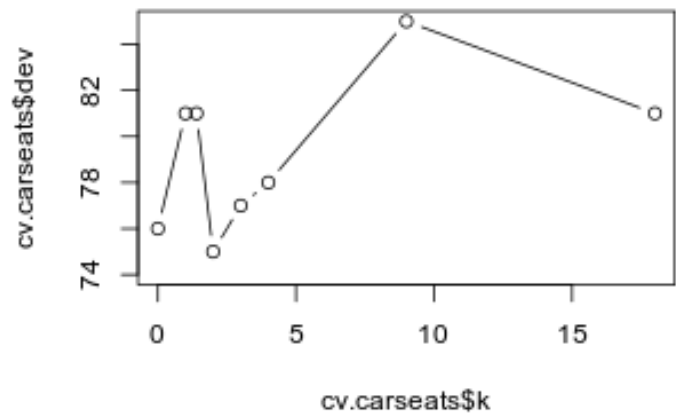
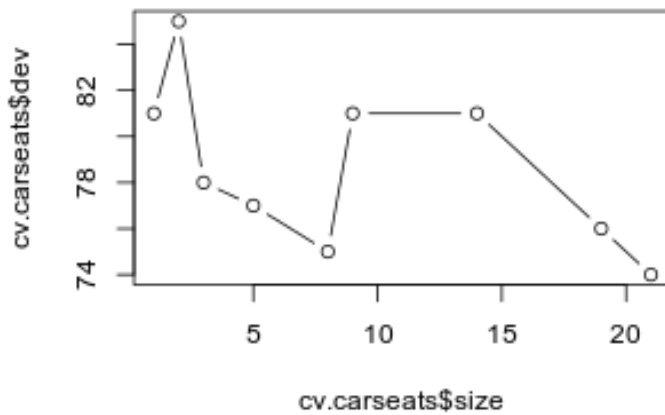
```
attr("class")
```

```
[1] "prune"          "tree.sequence"
```

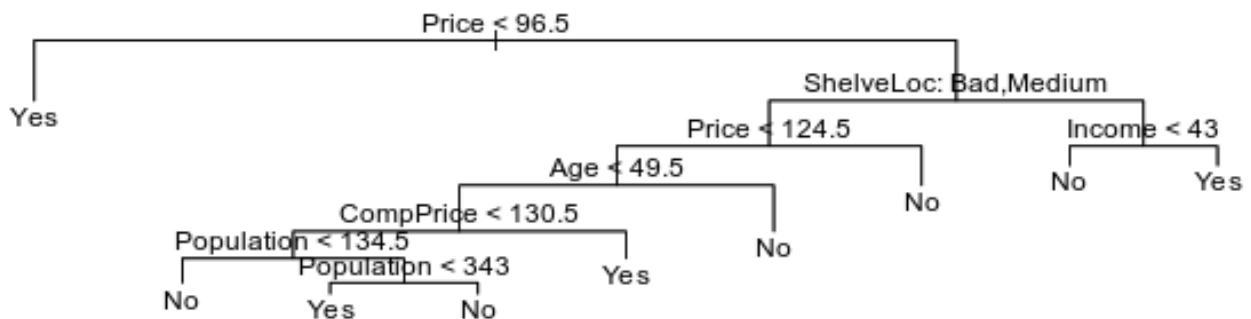
```

par(mfrow = c(1, 2))
plot(cv.carseats$size, cv.carseats$dev, type = "b")
plot(cv.carseats$k, cv.carseats$dev, type="b")

```



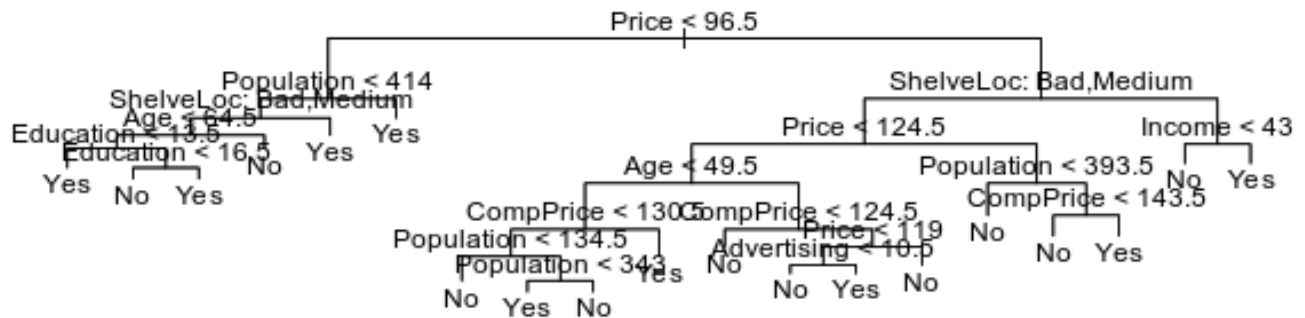
```
par(mfrow = c(1,1))
prune.carseats <- prune.misclass(tree.carseats, best = 9)
plot(prune.carseats)
text(prune.carseats, pretty = 0)
```



```
tree.pred <- predict(prune.carseats, carseats.test, type = "class")
table(tree.pred, high.test)
```

```
      high.test
tree.pred No Yes
No      97  25
Yes     20  58
```

```
prune.carseats <- prune.misclass(tree.carseats, best = 15)
plot(prune.carseats)
text(prune.carseats, pretty = 0)
```



```
tree.pred <- predict(prune.carseats, carseats.test, type = "class")
table(tree.pred, high.test)
```

```

      high.test
tree.pred No Yes
      No  102  30
      Yes   15  53
```

Regression Trees

```
boston <- as.data.table(Boston)

N <- nrow(boston)

set.seed(1)

train <- sample(1:N, N / 2)

tree.boston <- tree(medv ~ ., boston, subset = train)

summary(tree.boston)
```

Regression tree:

```
tree(formula = medv ~ ., data = boston, subset = train)
```

Variables actually used in tree construction:

```
[1] "rm" "lstat" "crim" "age"
```

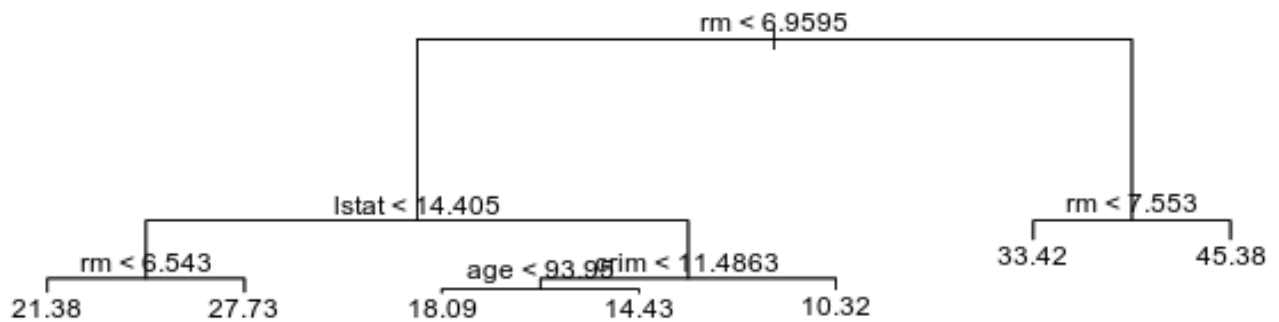
Number of terminal nodes: 7

Residual mean deviance: 10.38 = 2555 / 246


Distribution of residuals:

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|----------|---------|---------|--------|---------|---------|
| -10.1800 | -1.7770 | -0.1775 | 0.0000 | 1.9230 | 16.5800 |

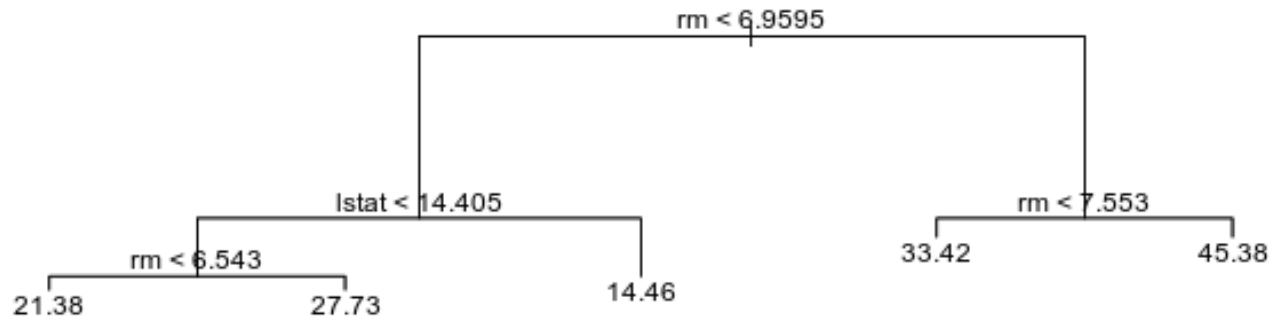
```
plot(tree.boston)
text(tree.boston, pretty = 0)
```



```
cv.boston <- cv.tree(tree.boston)
plot(cv.boston$size, cv.boston$dev, type = 'b')
```

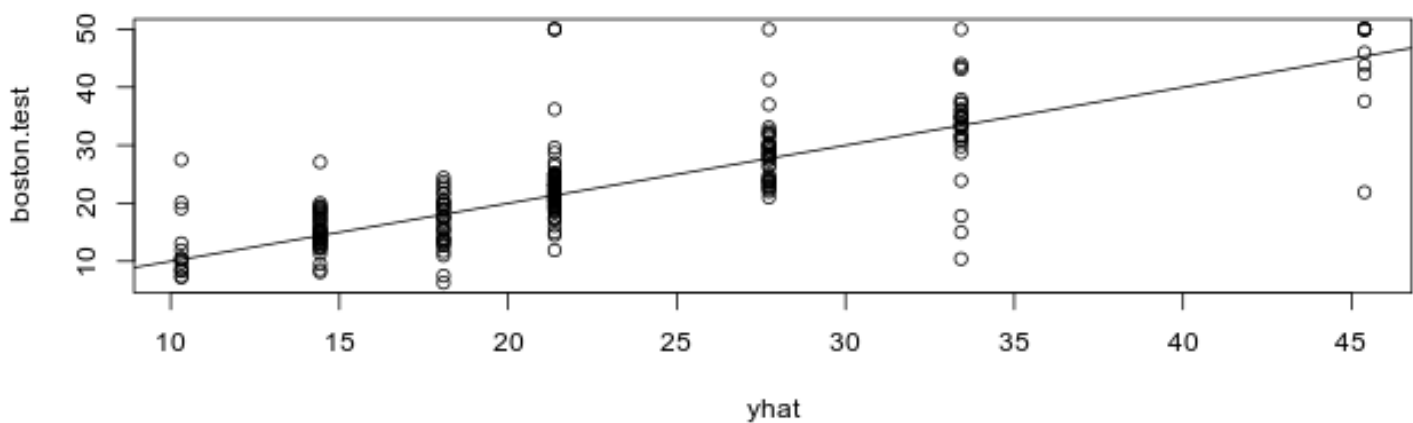


```
prune.boston <- prune.tree(tree.boston, best = 5)
plot(prune.boston)
text(prune.boston, pretty = 0)
```



```
yhat <- predict(tree.boston, newdata = boston[-train,])
boston.test <- boston[-train]$medv

plot(yhat, boston.test)
abline(0, 1)
```



```
mean((yhat - boston.test)^2)
```

```
[1] 35.28688
```

Bagging and Boosting

```
set.seed(1)
```

```
bag.boston <- randomForest(medv ~ ., data = boston, subset = train, mtry = 13, importance = T)
bag.boston
```

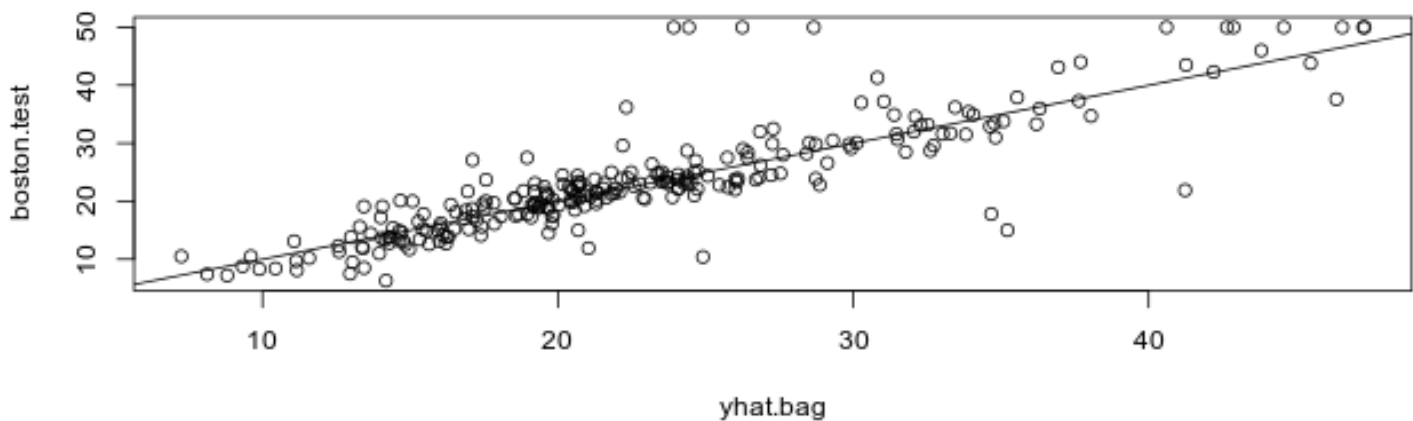
Call:

```
randomForest(formula = medv ~ ., data = boston, mtry = 13, importance = T,      subset = train
              Type of random forest: regression
              Number of trees: 500
```

```
No. of variables tried at each split: 13
```

```
Mean of squared residuals: 11.39601
% Var explained: 85.17
```

```
yhat.bag <- predict(bag.boston, newdata = boston[-train])
plot(yhat.bag, boston.test)
abline(0, 1)
```



```
mean((yhat.bag - boston.test)^2)
```

```
[1] 23.59273
```

```
bag.boston <- randomForest(medv ~ ., data = boston, subset = train, mtry = 13, ntree = 25)
yhat.bag <- predict(bag.boston, newdata = boston[-train,])
mean((yhat.bag - boston.test)^2)
```

```
[1] 23.66716
```

```
set.seed(1)
```

```
rf.boston <- randomForest(medv ~ ., data = boston, subset = train, mtry = 6, importance = T)
yhat.rf <- predict(rf.boston, newdata = boston[-train,])
mean((yhat.rf-boston.test)^2)
```


```
[1] 19.62021
```

```
importance(rf.boston)
```

| | %IncMSE | IncNodePurity |
|---------|-----------|---------------|
| crim | 16.697017 | 1076.08786 |
| zn | 3.625784 | 88.35342 |
| indus | 4.968621 | 609.53356 |
| chas | 1.061432 | 52.21793 |
| nox | 13.518179 | 709.87339 |
| rm | 32.343305 | 7857.65451 |
| age | 13.272498 | 612.21424 |
| dis | 9.032477 | 714.94674 |
| rad | 2.878434 | 95.80598 |
| tax | 9.118801 | 364.92479 |
| ptratio | 8.467062 | 823.93341 |
| black | 7.579482 | 275.62272 |
| lstat | 27.129817 | 6027.63740 |

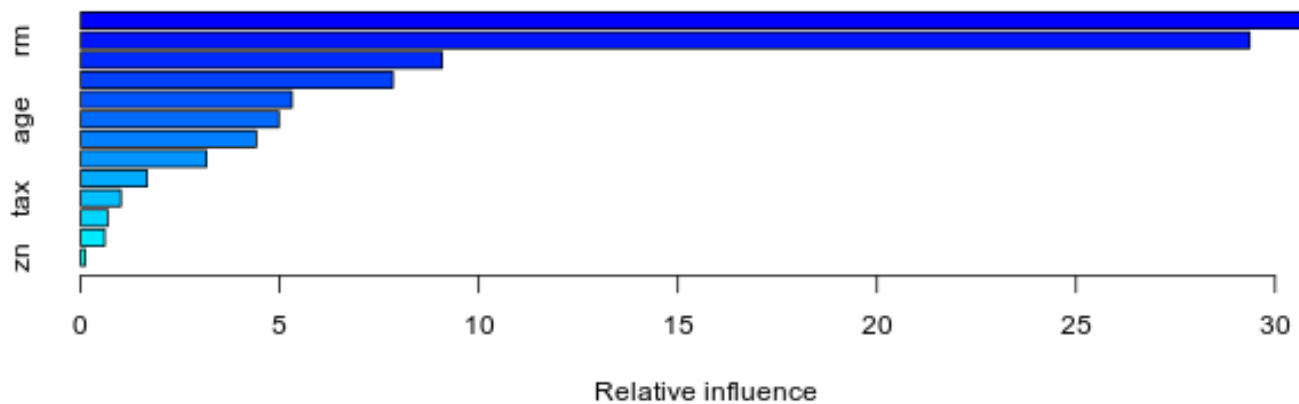
```
varImpPlot(rf.boston)
```

rf.boston



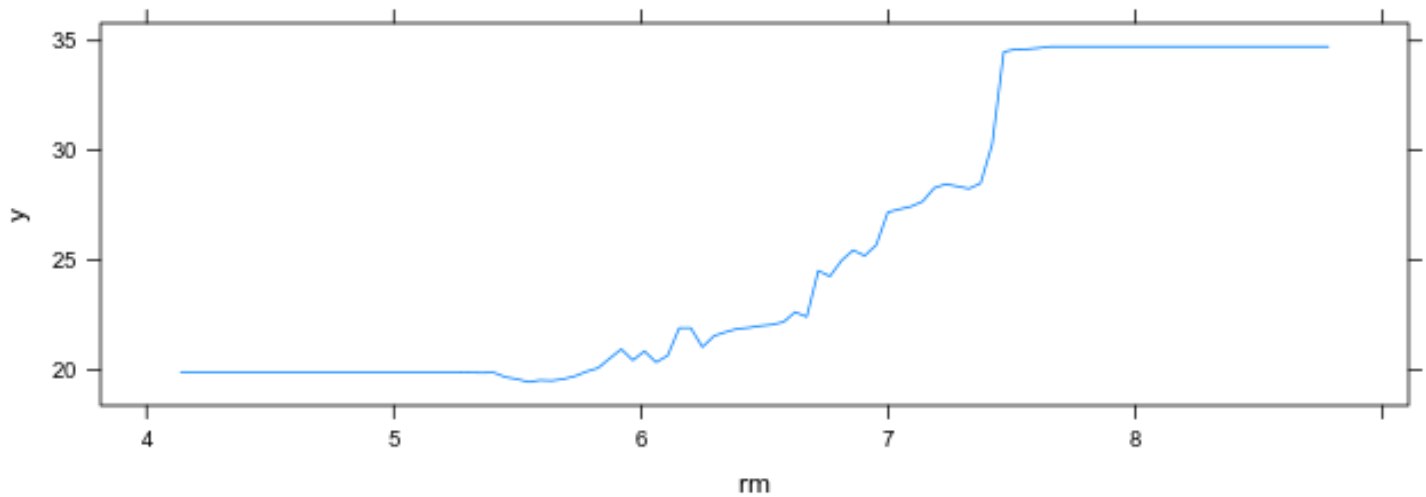
Boosting

```
boost.boston <- gbm(medv ~ ., data = boston[-train,], distribution = "gaussian", n.trees = 5000)
summary(boost.boston)
```

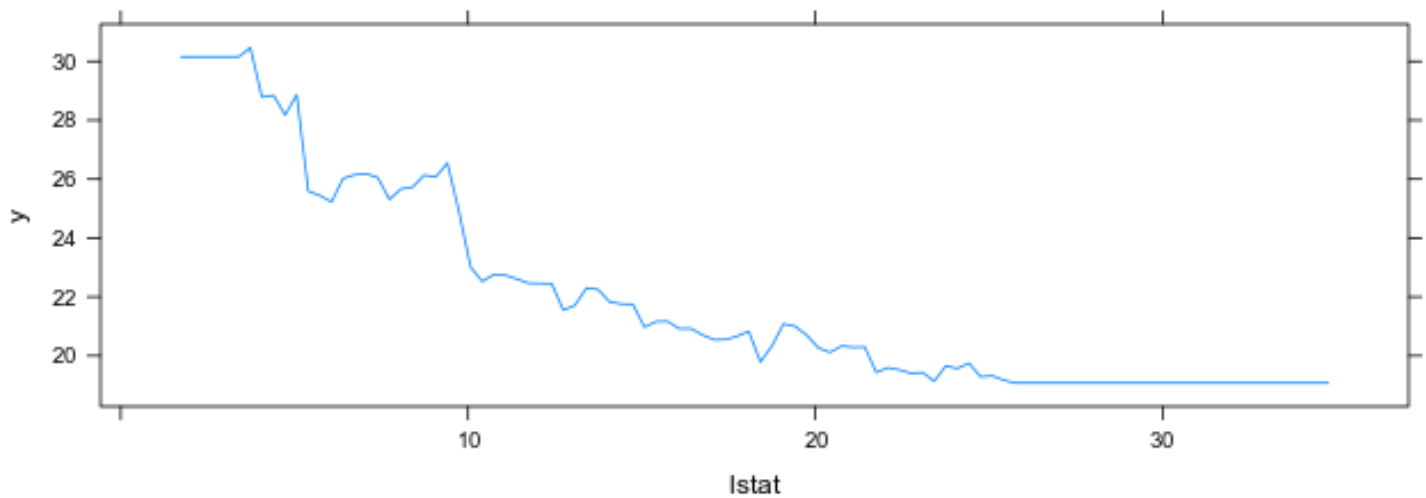


| | var | rel.inf |
|---------|---------|------------|
| lstat | lstat | 31.7768815 |
| rm | rm | 29.3618406 |
| dis | dis | 9.0801247 |
| crim | crim | 7.8431133 |
| nox | nox | 5.3077691 |
| age | age | 4.9799119 |
| black | black | 4.4118851 |
| ptratio | ptratio | 3.1614540 |
| indus | indus | 1.6642386 |
| tax | tax | 1.0107120 |
| chas | chas | 0.6850917 |
| rad | rad | 0.6031076 |
| zn | zn | 0.1138700 |

```
par(mfrow = c(1, 2))
plot(boost.boston, i = "rm")
```



```
plot(boost.boston, i = "lstat")
```



```
yhat.boost <- predict(boost.boston, newdata = boston[-train,], n.trees = 5000)
mean((yhat.boost - boston.test)^2)
```

```
[1] 2.242007e-06
```

Applied

In the lab, we applied random forests to the *Boston* data using `mtry = 6` and using `ntree = 25` and `ntree = 500`. Create a plot displaying the test error resulting from random forests on this data set for a more comprehensive range of values for `mtry` `ntree`.

```

N <- nrow(boston)

train <- sample(1:N, N * .7)

boston.train <- boston[train]
boston.test <- boston[!train]

ntrees <- seq(0, 500, 10)

results <- numeric(length(ntrees))

for(i in 1:length(ntrees))
{
  tree.count <- ntrees[i]

  rf <- randomForest(medv ~ ., data = boston.train, mtry = 6, n.trees = tree.count)

  pred <- predict(rf, newdata = boston.test)

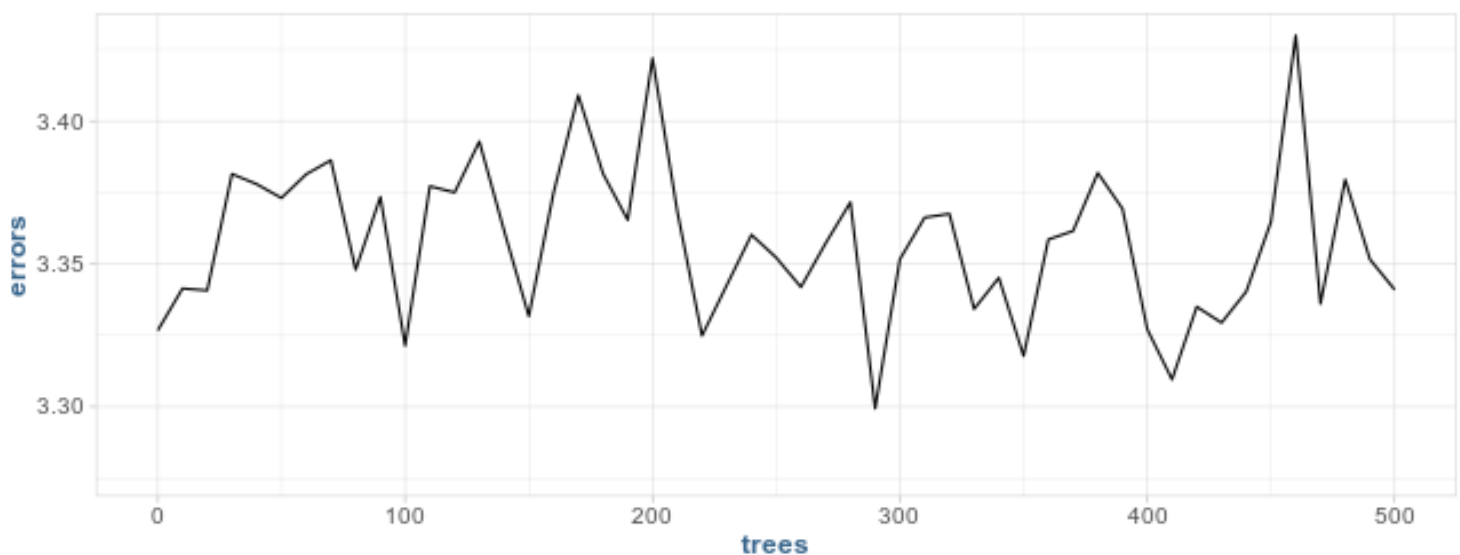
  results[i] <- sqrt(mean((pred - boston.test$medv)^2)) # store the rmse
}

lowest.error <- which.min(results)

rf.results <- data.table(trees = ntrees, errors = results)[, lowest := .I == lowest.error]

ggplot(rf.results, aes(trees, errors, fill = lowest)) +
  geom_line()

```



In the lab, a classification tree was applied to the Carseats data set after converting Sales into a qualitative response variable. Now we will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable.

a.) Split the data into a training set and a test set.

```
N <- nrow(carseats)

index <- sample(1:N, N * .7)

carseats.train <- carseats[index]
carseats.test <- carseats[!index]
```

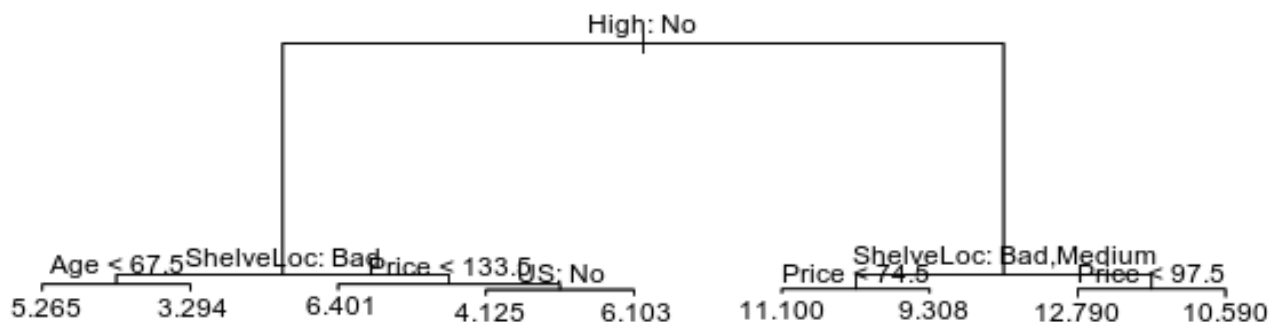
b.) Fit a regression tree to the training set. Plot the tree, and interpret the results. What test MSE do you obtain?

```
carseat.tree <- tree(Sales ~., data = carseats.train)
pred <- predict(carseat.tree, newdata = carseats.test)
mse <- mean((pred - carseats.test$Sales)^2)
```

```
mse
```

```
[1] 3.004485
```


```
par(mfrow = c(1,1))
plot(carseat.tree)
text(carseat.tree, pretty = 0)
```




c.) Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test MSE?

```
carseat.cv <- cv.tree(carseat.tree)

plot(carseat.cv$size, carseat.cv$dev, type = "b", col = "red")
```




```
plot(carseat.cv$size, carseat.cv$k, type = "b", col = "red")
```



```
prune.carseats <- prune.tree(carseat.tree, best = 5)
```

```
plot(prune.carseats)  
text(prune.carseats)
```



d.) Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important.

```

p <- ncol(carseats)

ntrees <- seq(0, 500, 25)
rf.error <- numeric(length(ntrees))

for(i in 1:length(ntrees))
{
  carseats.rf <- randomForest(Sales ~ ., data = carseats.train, mtry = p, n.trees = ntrees[i])

  pred <- predict(carseats.rf, newdata = carseats.test)

  rf.error[i] <- mean((pred - carseats.test$Sales)^2)
}

```

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

range

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

```
lowest.error <- which.min(rf.error)

results <- data.table(trees = ntrees, error = rf.error)[, lowest := .I == lowest.error]

ggplot(results, aes(trees, error, fill = lowest)) +
  geom_bar(stat = "identity") +
  labs(title = paste0("lowest error: #", ntrees[lowest.error]))
```



e.) Use random forest to analyze this data. What test MSE do you obtain? Use the importance() function to determine which variables are most important. Describe the effect of m , the number of variables considered at each split, on the error obtained.

```
p <- ncol(carseats)

mtry <- seq(1, p)

rf.error <- numeric(p)

for(i in 1:p)
{
  carseats.rf <- randomForest(Sales ~ ., data = carseats.train, mtry = mtry[i], n.trees = 250)

  pred <- predict(carseats.rf, newdata = carseats.test)

  rf.error[i] <- mean((pred - carseats.test$Sales)^2)
}
```


Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

```
lowest.error <- which.min(rf.error)
```

```
results <- data.table(mtry = mtry, error = rf.error, lowest = mtry == lowest.error)
```

```
ggplot(results, aes(mtry, error, fill = lowest)) +  
  geom_bar(stat = "identity")
```

