

Chapter 6

6.1

Explain the meaning of a .95 confidence interval.

A confidence interval is a statistical measure of the probability coverage a given lies within some interval based upon the observed data.

6.2

If the goal is to compute a .80, or .92, or a .98 confidence interval for μ when σ is known, and sampling is from a normal distribution, what values for c should be used in Equation (6.4)?

```
pretty_kable(
  data.table(
    Confidence = c(.8, .92, .98) ),
    Val := (1 - (1 - Confidence)/2) [, C := qnorm(Val)] ,
  "Confidence Values")
```

Table 1: Confidence Values

Confidence	Val	C
0.80	0.90	1.28
0.92	0.96	1.75
0.98	0.99	2.33

6.3

```
conf <- function(alpha) {
  c( Lower = (1 - alpha)/2, Upper = 1 - (1 - alpha)/2)
}
```

Assuming random sampling is from a normal distribution with standard deviation $\sigma = 5$, if the sample mean is $\bar{X} = 45$ based on $n = 25$ participants, what is the 0.95 confidence interval for μ ?

```
qnorm(conf(.95), mean = 45, sd = 5/sqrt(25))
```

```
Lower    Upper
43.04004 46.95996
```

6.4

Repeat the previous example, only compute a .99 confidence interval instead.

```
qnorm(conf(.99), mean = 45, sd = 5/sqrt(25))
```

Lower	Upper
42.42417	47.57583

6.5

A manufacturer claims that their light bulbs have an average life span that follows a normal distribution with $\mu = 1,200$ hours and a standard deviation of $\sigma = 25$. If you randomly test 36 light bulbs and find that their average life span is $\bar{X} = 1,150$, does a .95 confidence interval for μ suggest that the claim $\mu = 1,200$ is reasonable?

```
qnorm(conf(.95), mean = 1150, sd = 25/sqrt(36))
```

Lower	Upper
1141.833	1158.167

No, 1,200 is outside the bounds of a 95% confidence interval.

6.6

Compute a .95 confidence interval for the mean, assuming normality, for the following situations:

a.) $n = 12, \sigma = 22, \bar{X} = 65$

```
qnorm(conf(.95), mean = 65, sd = 22/sqrt(12))
```

Lower	Upper
52.55256	77.44744

b.) $n = 22, \sigma = 10, \bar{X} = 185$

```
qnorm(conf(.95), mean = 185, sd = 10/sqrt(22))
```

Lower	Upper
180.8213	189.1787

c.) $n = 50, \sigma = 30, \bar{X} = 19$

```
qnorm(conf(.95), mean = 19, sd = 30/sqrt(50))
```

Lower	Upper
10.68458	27.31542

6.7

What happens to the length of a confidence interval for the mean of a normal distribution when the sample size is doubled? In particular, what is the ratio of the lengths?

What is the ratio of the lengths if the sample size is quadrupled?

Answer:

In general, For some x ,

$$CI = \mu \pm C \frac{\sigma}{x\sqrt{n}}$$

As x increases, the size of the confidence intervals decrease (closer approximations).

6.8

The length of a bolt made by a machine parts company has a normal distribution with standard deviation σ equal to 0.01 mm. The lengths of four randomly selected bolts are as follows:

20.01, 19.88, 20.00, 19.99

a.) Compute a 0.95 confidence interval for the mean.

```
x <- c(20.01, 19.88, 20.00, 19.99)
n <- length(x)
xbar <- mean(x)

qnorm(conf(.95), mean = xbar, sd = 0.01/sqrt(n))
```

```
Lower    Upper
19.9602  19.9798
```

b.) Specifications require a mean lengths μ of 20.00 mm for the population of bolts. Do the data indicate that this specification is being met?

No, the 95% confidence interval for the mean does not include 20.

c.) Given that the 0.95 confidence interval contains the value 20, why might it be inappropriate to conclude that the specification is being met?

Just because the confidence interval includes some value, that does not mean that the population parameter is that value.

6.9

The weight of trout sold at a trout farm has a standard deviation of 0.25. Based on a sample of 10 trout, the average weight is 2.10 lbs.

Assume normality and Compute a .99 confidence interval for the mean.

```
qnorm(conf(.99), mean = 2.10, sd = .25/sqrt(10))
```

Lower	Upper
1.896363	2.303637

6.10

The average bounce of 45 randomly selected tennis balls is found to be $\bar{X} = 1.70$.

Assuming that the standard deviation of the bounce is .30, compute a 0.90 confidence interval for the average bounce assuming normality.

```
qnorm(conf(.90), mean = 1.70, sd = .30/sqrt(45))
```

Lower	Upper
1.62644	1.77356

6.11

Assuming that the degrees of freedom are 20, find the value t for which:

a.) $P(T \leq t) = 0.995$

```
qt(.995, df = 20)
```

```
[1] 2.84534
```

b.) $P(T \geq t) = 0.025$

```
1 - qt(.025, df = 20)
```

```
[1] 3.085963
```

c.) $P(-t \leq T \leq t) = 0.90$

```
qt((1 + .90) / 2, df = 20)
```

```
[1] 1.724718
```

6.12

Compute a 0.95 confidence interval if

a.) $n = 10, \bar{X} = 26, s = 9$

```
26 + qt(conf(.95), df = 9) * 9/sqrt(10)
```

```

Lower    Upper
19.56179 32.43821

```

b.) $n = 18, \bar{X} = 132, s = 20$

```
132 + qt(conf(.95), df = 17) * 20/sqrt(18)
```

```

Lower    Upper
122.0542 141.9458

```

c.) $n = 25, \bar{X} = 52, s = 12$

```
52 + qt(conf(.95), df = 24) * 12/sqrt(25)
```

```

Lower    Upper
47.04664 56.95336

```

6.13

Compute a 0.99 confidence interval if

a.) $n = 10, \bar{X} = 26, s = 9$

```
26 + qt(conf(.99), df = 9) * 9/sqrt(10)
```

```

Lower    Upper
16.75081 35.24919

```

b.) $n = 18, \bar{X} = 132, s = 20$

```
132 + qt(conf(.99), df = 17) * 20/sqrt(18)
```

```

Lower    Upper
118.3376 145.6624

```

c.) $n = 25, \bar{X} = 52, s = 12$

```
52 + qt(conf(.99), df = 24) * 12/sqrt(25)
```

```

Lower    Upper
45.28735 58.71265

```

6.14

For a study on self-awareness, the observed values for one of the groups were:

77, 87, 88, 114, 151, 210, 219, 246, 253, 262, 296, 299, 306, 376, 428, 515, 666, 1310, 2611.

Compute a .95 confidence interval for the mean assuming normality.

```
x <- c(77,87,88,114,151,210,219,246,253,262,296,299,306,376,428,515, 666,1310,2611)
xbar <- mean(x)
n <- length(x)

xbar + qt(conf(.95), df = n - 1) * sd(x)/sqrt(n)
```

```
      Lower      Upper
161.5030 734.7075
```

Compare the result to the bootstrap-t confidence interval.

```
winsorize <- function( x, trim = .2) {
  n <- length(x)

  o <- sort(x)
  g <- floor(trim*n)

  o[1:(g+1)] <- o[(g+1)]
  o[(n-g):n] <- o[n-g]

  o
}

trimse<-function(x,tr=.2,na.rm=FALSE){
  if(na.rm)x<-x[!is.na(x)]
  trimse <- sqrt(var(winsorize(x,tr)))/((1-2*tr)*sqrt(length(x)))
  trimse
}

nboot <- 2000
alpha <- .95

sims <- matrix(sample(x, size = x*nboot, replace =T), nrow = nboot)
data <- sims - xbar

top <- apply(data, 1, mean)
bot <- apply(data, 1, trimse)

tval <- top/bot
tval <- sort(tval)

icrit <- round((1-alpha)*nboot)

c(Lower = mean(x)+tval[icrit]*trimse(x),
  Upper = mean(x)-tval[icrit]*trimse(x))
```

Lower	Upper
162.7704	733.4401

Why do they differ?

Student's T gives a shorter confidence interval, however the probability coverage suggest less than 95% coverage.

6.15

Rats are subjected to a drug that might affect aggression. Suppose that for a random sample of rats, measures of aggression are found to be:

5, 12, 23, 24, 18, 9, 18, 11, 36, 15

Compute a 0.95 confidence interval for the mean assuming the scores are from a normal distribution.

```
x <- c(5, 12, 23, 24, 18, 9, 18, 11, 36, 15)

xbar <- mean(x); n <- length(x)

xbar + qt(conf(.95), df = n - 1) * sd(x)/sqrt(n)
```

Lower	Upper
10.69766	23.50234

6.16

Suppose $M = 34$ and the McKean-Schrader estimate of the standard error of M is $S_M = 3$.

Compute a .95 confidence interval for the population median.

```
34 + qnorm(conf(.95)) * 3
```

Lower	Upper
28.12011	39.87989

6.17

For the data in E14, the McKean-Schrader estimate of the standard error of M is $S_M = 77.8$ and the sample median is 262.

Compute a .99 confidence interval for the population median.

```
x <- c(77,87,88,114,151,210,219,246,253,262,296,299,306,376,428,515, 666,1310,2611)
M <- median(x); n <- length(x)

M + qnorm(conf(.99)) * 77.8
```

Lower	Upper
61.60048	462.39952

6.18

If $n = 10$ and a confidence interval for the median is computed using (X_k, X_{n-k}) as described in 6.3.4, what is the probability coverage if $k = 2$?

```
1 - 2*pbinom(1, size = 10, prob = .5)
```

```
[1] 0.9785156
```

6.19

From the previous, what is the probability coverage when $n = 15$ and $k = 4$?

```
1 - 2*pbinom(3, size = 15, prob = .5)
```

```
[1] 0.9648438
```

6.20

For the data in E14, if you use (88, 515) as a confidence interval for the median, what is the probability this interval contains the population median?

```
x <- c(77, 87, 88, 114, 151, 210, 219, 246, 253, 262, 296, 299, 306, 376, 428, 515, 666, 1310,
```

```
k <- 3
```

```
1 - 2*pbinom(k, size = n - k - 1, prob = .5)
```

```
[1] 0.9648438
```

6.21

You observe the following successes and failures:

```
1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0
```

Compute a .95 confidence interval for p using E6.15, as well as the Agresti-Coull method.

```
x <- c(1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)
```

```
n <- length(x)
```

```
alpha <- 0.05
```



```
phat <- sum(x) / n

phat + qnorm(conf(1 - alpha)) * sqrt(phat*(1 - phat)/n)
```

```
      Lower      Upper
0.09477411 0.57189255
```

```
c <- qnorm(1 - alpha/2)
nhat <- n + c^2
xhat <- sum(x) + c^2/2
phat <- xhat / nhat

c(Lower = phat - c * sqrt(phat*(1 - phat)/nhat),
  Upper = phat + c * sqrt(phat*(1 - phat)/nhat))
```

```
      Lower      Upper
0.1496413 0.5849864
```

6.22

Given the following results for a sample from a binomial distribution, compute the squared standard error of \hat{p} .

a.) $n = 25$, $X = 5$

```
n <- 25; success <- 5
phat <- success/n
```

```
phat*(1 - phat)/n
```

```
[1] 0.0064
```

b.) $n = 48$, $X = 12$

```
n <- 48; success <- 12
phat <- success/n
```

```
phat*(1 - phat)/n
```

```
[1] 0.00390625
```

c.) $n = 100$, $X = 80$

```
n <- 100; success <- 80
phat <- success/n
```

```
phat*(1 - phat)/n
```

```
[1] 0.0016
```

d.) $n = 300$, $X = 160$

```
n <- 300; success <- 160
phat <- success/n

phat*(1 - phat)/n

[1] 0.0008296296
```

6.23

Among 100 randomly sampled adults, 10 were found to be unemployed.

```
n <- 100; success <- 10
phat <- success/n

phat + qnorm(conf(.95))*sqrt(phat*(1-phat)/n)
```

Lower	Upper
0.04120108	0.15879892

Compute a 0.95 confidence interval for the percentage of adults unemployed using E6.15.

Compare this result to the Agresti-Coull confidence interval.

```
alpha <- .05

c <- qnorm(1 - alpha/2)
nhathat <- n + c^2
xhat <- success + c^2/2
phat <- xhat / nhathat

c(Lower = phat - c * sqrt(phat*(1 - phat)/nhathat),
  Upper = phat + c * sqrt(phat*(1 - phat)/nhathat))
```

Lower	Upper
0.05348475	0.17611005

6.24

A sample of 1,000 fishes was obtained from a lake. It was found that 290 were members of the bass family.

Compute a 0.95 confidence interval for the percentage of bass in the lake using E6.15.

```
n <- 1000; success <- 290
phat <- success/n
```

```
phat + qnorm(conf(.95))*sqrt(phat*(1-phat)/n)
```

Lower	Upper
0.2618761	0.3181239

6.25

Among a random sample of 1,000 adults, 60 reported never having any legal problems. Use Equation 6.15 to compute the .95 confidence interval for the percentage of adults who have never had legal problems.

```
n <- 1000; success <- 60
```

```
phat <- success/n
phat + qnorm(conf(.95))*sqrt(phat*(1-phat)/n)
```

Lower	Upper
0.04528068	0.07471932

6.26

Among a sample of 600 items, only one was found to be defective. Explain why E6.15 might be unsatisfactory when computing a confidence interval for the probability that an item is defective.

When the probability is very low (or high) the equation in 6.15 performs poorly.

6.27

In the previous exercise, compute a 0.90 confidence interval for the probability that an item is defective.

```
n <- 600; success <- 1
phat <- success / n
```

```
alpha <- .10
```

```
round(c(Lower = 1 - (1 - alpha/2)^(success/n),
        Upper = 1 - (alpha/2)^(success/n)), 5)
```

Lower	Upper
0.00009	0.00498

6.28

One-fourth of 300 persons in a large city stated that they are opposed to a certain political issue favored by the mayor.

Using E6.15, calculate a .99 confidence interval for the fraction of individuals opposed to this issue.

```
n <- 300; success <- n * .25
phat <- success/n

phat + qnorm(conf(.99)) * sqrt(phat*(1-phat)/n)

      Lower      Upper
0.1856043 0.3143957
```

6.29

Consider a method for detecting a type of cancer that is highly accurate but very expensive and invasive, resulting in a great deal of discomfort for the patient. An alternative test to detect this type of cancer has been developed, and it is of interest to know the probability of a false-negative indication, meaning that the test fails to detect cancer when it is present. The test is conducted on 250 patients known to have cancer, and 5 tests fail to show its presence. Use equation 6.15 to determine a 0.95 confidence interval for the probability of a false-negative indication.

```
n <- 250; success <- 5
phat <- success/n

phat + qnorm(conf(.95)) * sqrt(phat*(1-phat)/n)

      Lower      Upper
0.002645739 0.037354261
```

6.30

In the previous exercise, imagine that 0 false-negative indications were found.

Determine a .99 confidence interval for the probability of a false-negative indication.

```
n <- 250;

alpha <- 0.01

phat <- 1 - alpha^(1/250)

phat

[1] 0.01825206

round(c(Lower = 1 - (1 - alpha/2)^(success/n),
      Upper = 1 - (alpha/2)^(success/n)), 5)

      Lower      Upper
0.00010 0.10055
```

6.31

A cosmetic company found that 180 of 1,000 randomly selected women in New York city have seen the company's latest television ad. Use E6.15 to compute a .95 confidence interval for the percentage of women in New York city who have seen the ad.

```
n <- 1000; success <- 180
phat <- success/n

phat + qnorm(conf(.95)) * sqrt(phat*(1-phat)/n)
```

Lower	Upper
0.1561883	0.2038117

6.32

Describe in general terms how nonnormality can affect Student's T distribution.

The actual probability coverage can be substantially over or understated depending on the skewness.

6.33**6.34**

Listed here are the average LSAT scores for the 1973 entering class of 15 American law schools.

545,555,558,572,575,576,578,580, 594,605,635,651,653,661,666


The .95 confidence interval for the population mean μ is (577.1, 623.4).

Use a boxplot to verify that there are not outliers and that the data appear have a skewed distribution.

```
x <- c(545,555,558,572,575,576,578,580,594,605,635,651,653,661,666)

data <- data.table(values = x)

ggplot(data, aes(y = values, group = "")) +
  geom_boxplot(fill = "cornflowerblue") +
  geom_hline(yintercept = 577.1, color = "darkred") +
  geom_hline(yintercept = 623.4, color = "darkred")
```



6.35

Compute a 0.95 confidence interval for the 20% trimmed mean. if:

a.) $n = 24$, $s_w^2 = 12$, $\bar{X}_t = 52$

```
xbar <- 52; wt <- 12; n <- 24
```

```
g <- floor(n * .2)
```

```
xbar + qt(conf(.95), df = n - 2*g - 1) * sqrt(wt) / (.6*sqrt(n))
```

Lower	Upper
49.48806	54.51194

b.) $n = 36$, $s_w^2 = 30$, $\bar{X}_t = 10$

```
xbar <- 10; wt <- 30; n <- 36
```

```
g <- floor(n * .2)
```

```
xbar + qt(conf(.95), df = n - 2*g - 1) * sqrt(wt) / (.6*sqrt(n))
```

Lower	Upper
6.835968	13.164032

c.) $n = 12$, $s_w^2 = 9$, $\bar{X}_t = 16$

```
xbar <- 16; wt <- 9; n <- 12

g <- floor(n * .2)

xbar + qt(conf(.95), df = n - 2*g - 1) * sqrt(wt) / (.6*sqrt(n))

      Lower      Upper
12.58696 19.41304
```

6.36

Repeat with .99 confidence.

a.) $n = 24, s_w^2 = 12, \bar{X}_t = 52$

```
xbar <- 52; wt <- 12; n <- 24

g <- floor(n * .2)

xbar + qt(conf(.99), df = n - 2*g - 1) * sqrt(wt) / (.6*sqrt(n))

      Lower      Upper
48.52727 55.47273
```

b.) $n = 36, s_w^2 = 30, \bar{X}_t = 10$

```
xbar <- 10; wt <- 30; n <- 36

g <- floor(n * .2)

xbar + qt(conf(.99), df = n - 2*g - 1) * sqrt(wt) / (.6*sqrt(n))

      Lower      Upper
5.692224 14.307776
```

c.) $n = 12, s_w^2 = 9, \bar{X}_t = 16$

```
xbar <- 16; wt <- 9; n <- 12

g <- floor(n * .2)

xbar + qt(conf(.99), df = n - 2*g - 1) * sqrt(wt) / (.6*sqrt(n))

      Lower      Upper
10.94893 21.05107
```

6.37

The last example at the end of Section 6.3.1 used data from a study of self-awareness. In another portion of the study, a group of participants had the follow values.

59, 106, 174, 207, 219, 237, 313, 365, 458, 497, 515, 529, 557, 615, 625, 645, 973, 1065, 3215

Compute a .95 confidence interval for both the population mean and 20% trimmed mean.

```
x <- c(59, 106, 174, 207, 219, 237, 313, 365, 458, 497, 515, 529, 557, 615, 625,
645, 973, 1065, 3215)
```

```
xbar <- mean(x); n <- length(x)
```

```
xbar + qt(conf(.95), df = n - 1) * sqrt(sd(x) / n)
```

```
      Lower      Upper
585.9862 611.2770
```

```
g <- floor(n * .2)
```

```
xt <- x[g:n-g]
```

```
xtbar <- mean(xt); nt <- length(xt)
```

```
xtbar + qt(conf(.95), df = nt - 1) * sd(xt) / sqrt(nt)
```

```
      Lower      Upper
277.9201 487.2049
```

6.38

The ideal method for computing a confidence interval for the mean would have a shorter length than any other method that might be used. Explain why such a method does not exist even when sampling from a perfectly symmetric distribution.

Because outliers.

6.39

For the values:

5, 60, 43, 56, 32, 43, 47, 79, 39, 41

compute a .95 confidence interval for the 20% trimmed mean and compare the results to the .95 confidence interval for the mean.


```
x <- c(5, 60, 43, 56, 32, 43, 47, 79, 39, 41)

xbar <- mean(x); n <- length(x)
xbar + qt(conf(.95), df = n - 1) * sqrt(sd(x)/n)
```

```
      Lower      Upper
41.36199 47.63801
```

```
g <- floor(n * .2)
```

```
x <- sort(x)
x <- x[g:n-g]
```

```
xbar <- mean(x); n <- length(x)
xbar + qt(conf(.95), df = n - 1) * sqrt(sd(x)/n)
```

```
      Lower      Upper
35.00468 41.49532
```

6.40

In the previous exercise, the confidence interval for the 20% trimmed mean is shorter than the confidence interval for the mean.

```
outbox(x)
```

```
$out.val
[1] 5
```

```
$out.id
[1] 1
```

```
$keep
[1] 2 3 4 5 6 7 8
```

```
$n
[1] 8
```

```
$n.out
[1] 1
```

```
$cl
[1] 19.29167
```

```
$cu
[1] 60.95833
```

Explain why this is not surprising.

6.41

The following code for estimating the distribution of Student's T when data are sampled from a skewed, light-tailed distribution that has a population mean $\mu = 1.649$:

```
t.vals <- NULL
for(study in 1:1000) {
  x <- rlnorm(25)
  suppressMessages({
    t.vals[study] <- trimciv2(x, tr=0, null.value=1.649)$test.stat})
}
```

6.42

Modify the previous code to use actual data from a study to compute a confidence interval.

```
t.vals=NULL
for(i in 1:1000){
  z=sample(x,25,replace=T)
  t.vals[i]=trimci(z,tr=0,null.value=mean(x),pr=F)$test.stat
}

ggplot(data.table(values = x), aes(values)) +
  geom_histogram(aes(y = ..density.., fill = ..count..), bins = 30) +
  geom_density(fill = "darkorange", alpha = .3)
```

