

Chapter 8

8.1

For the following data, use R to verify that the least squares regression line is $\hat{Y} = 1.8X - 8.5$

X: 5, 8, 9, 7, 14 Y: 3, 1, 6, 7, 19

```
dat <- data.table(
  X = c(5, 8, 9, 7, 14),
  Y = c(3, 1, 6, 7, 19))

lm(Y ~ X, data = dat)
```

Call:

```
lm(formula = Y ~ X, data = dat)
```

Coefficients:

(Intercept)	X
-8.478	1.823

Also verify that the Theil-Sen estimator, the slope is estimated to be 1.746 and the intercept is estimated to be -7.968.

```
tsreg(dat$X, dat$Y)$coef
```

Intercept

```
-5.730159  1.746032
```

8.2

Using the R function *lsfit*, compute the residuals using the data in E1,

Verify that if you square and sum the residuals, you get 46.585.

```
res <- lsfit(dat$X, dat$Y)$resid

sum(res^2)
```

```
[1] 46.58407
```

8.3

Verify that for the data in E1, if you use $\hat{Y} = 2X - 9$, the sum of the squared residuals is greater than 46.584.

```
Yhat <- 2*dat$X - 9
res <- sum( (dat$Y - Yhat)^2 )
```

```
stopifnot(res > 46.583)
res
```

```
[1] 53
```

Why would you expect a value greater than 46.584?

The x coefficient increased.

8.4

Suppose that based on $n = 25$ values, $s_x^2 = 12$ and $\sum (X_i - \bar{X})(Y_i - \bar{Y}) = 144$.

What is the slope of the least squares regression?

$A = 144, C = (n - 1)s_x^2 = 288, b_1 = A/C = 144/288 = .5$

8.5

The following table reports breast cancer rates plus levels of solar radiation (in calories per day) for various cities in the United States. The data are stored in the file `cancer_rate_dat.txt`.

```
dat <- data.table::fread(paste0(data.dir, "cancer_rate_dat.txt"), fill = T, sep = "&")
```

```
dat
```

	City	Rate	calories
1:	New York	32.75	300
2:	Chicago	30.75	275
3:	Pittsburgh	28.00	280
4:	Seattle	27.25	270
5:	Boston	30.75	305
6:	Cleveland	31.00	335
7:	Columbus	29.00	340
8:	Indianapolis	26.50	342
9:	New Orleans	27.00	348
10:	Nashville	23.50	354
11:	Washington, DC	31.20	357
12:	Salt Lake City	22.70	394
13:	Omaha	27.00	380
14:	San Diego	25.80	383
15:	Atlanta	27.00	397
16:	Los Angeles	27.80	450
17:	Miami	23.50	453
18:	Fort Worth	21.50	446
19:	Tampa	21.00	456

```



20:    Albuquerque 22.50    513
21:      Las Vegas 21.50    510
22:      Honolulu 20.60    520
23:      El Paso  22.80    535
24:      Phoenix  21.00    520
      City Rate calories


```

Fit a OLS regression to predict cancer rates and comment on what this suggests.

```
fit <- lm(Rate ~ calories, data = dat)
```

```
plot(fit)
```





8.6

For the following data, use R to compute the least squares regression line for predicting GPA given SAT.

SAT: 500, 530, 590, 660, 610, 700, 570, 640 GPA: 2.3, 3.1, 2.6, 3.0, 2.4, 3.3, 2.6, 3.5

```
dat <- data.table(
  SAT = c(500, 530, 590, 660, 610, 700, 570, 640),
  GPA = c(2.3, 3.1, 2.6, 3.0, 2.4, 3.3, 2.6, 3.5)
)

fit <- lm(GPA ~ SAT, data = dat)
coef(fit)
```

```
(Intercept)      SAT
0.484615385 0.003942308
```

8.7

Compute the residuals for the data used in the previous problem and verify that sum to zero.

```
round(sum((dat$GPA - fit$fitted.values)), 5)
```

```
[1] 0
```

8.8

For the following data, use R to compute the least squares regression line for predicting Y from X.

X: 40, 41, 42, 43, 44, 45, 46 Y: 1.62, 1.63, 1.90, 2.64, 2.05, 2.13, 1.94

```
dat <- data.table(
  X = c(40, 41, 42, 43, 44, 45, 46),
  Y = c(1.62, 1.63, 1.90, 2.64, 2.05, 2.13, 1.94)
)

summary(fit <- lm(Y ~ X, data = dat))
```

Call:

```
lm(formula = Y ~ X, data = dat)
```

Residuals:

1	2	3	4	5	6	7
-0.141071	-0.206429	-0.011786	0.652857	-0.012500	-0.007857	-0.273214

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.25321	2.73157	-0.459	0.666
X	0.07536	0.06346	1.188	0.288

Residual standard error: 0.3358 on 5 degrees of freedom

Multiple R-squared: 0.22, Adjusted R-squared: 0.064

F-statistic: 1.41 on 1 and 5 DF, p-value: 0.2883

```
ggplot(dat, aes(X, Y)) +
  geom_point() +
  geom_smooth(method = "lm")
```




8.9

In exercise 5, what would be the least squares estimate of the cancer rate given a solar radiation of 600?

```
dat <- data.table::fread(paste0(data.dir, "cancer_rate_dat.txt"), fill = T, sep = "&")

ggplot(dat, aes(calories, Rate)) +
  geom_point() +
  geom_smooth(method = "lm")
```



```
fit <- lm(Rate ~ calories, data = dat)
```

```
coef(fit)
```

```
(Intercept)    calories  
39.99094634 -0.03565283
```

```
39.99 -0.037*600
```

```
[1] 17.79
```

Why might this be unreasonable?

Because 600 is outside of the bounds of seen values (extrapolation).

8.10

Maximal oxygen uptake (MOU) is a measure of an individual's physical fitness. You want to know how MOU is related to how fast someone can run a mile. Suppose you randomly sample six athletes and get:

MOU (kl/kg): 63.3, 60.1, 53.6, 58.8, 67.5, 62.5 Time (seconds): 241.5, 249.8, 246.1, 232.4, 237.2, 238.4

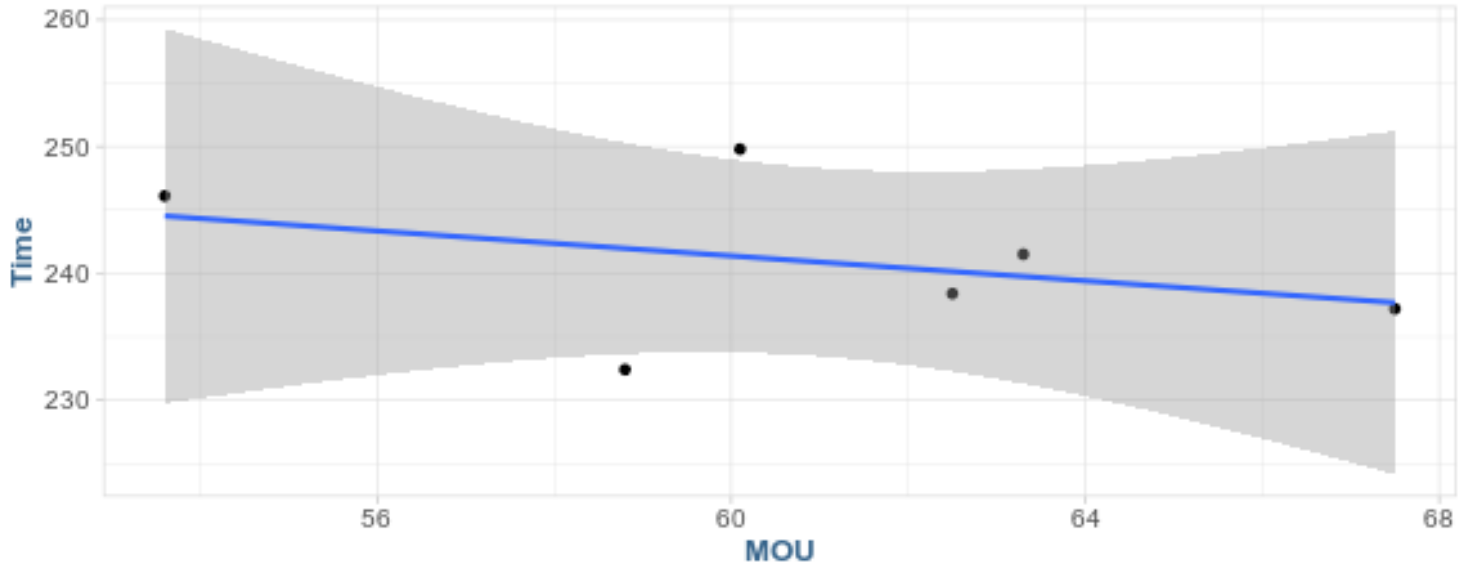
Compute the least squares regression line and comment on what the results suggest.

```
dat <- data.table(  
  MOU = c(63.3, 60.1, 53.6, 58.8, 67.5, 62.5),  
  Time = c(241.5, 249.8, 246.1, 232.4, 237.2, 238.4)  
)
```

```
fit <- lm(Time ~ MOU, data = dat)
```

```
ggplot(dat, aes(MOU, Time)) +  
  geom_point() +  
  geom_smooth(method = "lm") +  
  labs(title = "Time vs. MOU")
```

Time vs. MOU



Generally, time decreases as MOU increases.

8.11

Verify that for the following pairs of points, the least squares regression line has a slope of zero. Plot the points and comment on the assumption that the regression line is straight.

X: 1, 2, 3, 4, 5, 6 Y: 1, 4, 7, 7, 4, 1

```
dat <- data.table(
  X = c(1, 2, 3, 4, 5, 6),
  Y = c(1, 4, 7, 7, 4, 1)
)
```

```
fit <- lm(Y ~ X, data = dat)
```

```
coef(fit)
```

```
(Intercept)          X
4.000000e+00 -5.838669e-16
```

8.12

Repeat the last exercise, only for points:

X: 1, 2, 3, 4, 5, 6 Y: 4, 5, 6, 7, 8, 2


```
dat <- data.table(
  X = c(1, 2, 3, 4, 5, 6),
  Y = c(4, 5, 6, 7, 8, 2)
)

fit <- lm(Y ~ X, data = dat)

coef(fit)
```

```
(Intercept)          X
5.333333e+00 -6.369458e-16
```

8.13

Vitamin A is required for good health. However, one bite of polar bear liver results in death because it contains a high concentration of vitamin A. Comment on this fact in terms of extrapolation.


Vitamin A as a function of health is a bounded variable. Predicting a response when the independent variables are far from what has been observed can lead to undefined results (death).

8.14

Socket et al. (1987) report data related to patterns of residual insulin secretion in children. A portion of the study was concerned with whether age can be used to predict the logarithm of C-peptide concentrations at diagnosis. The observed values are (data file), Replicate the LOESS smoothed curve in fig 8.4.1.

```
dat <- data.table::fread(paste0(data.dir, "diabetes_sockett_dat.txt"), fill = T)

ggplot(dat, aes(age, pep)) +
  geom_point() +
  geom_smooth(method = "loess")
```



8.15

For the data in the last exercise, use R to verify that a least squares regression line using only X values (age) less than or equal to 7 yields a p-value equal to 0.026 when using the R function `olsch4`. Also verify that the p-value, when using the Theil-Sen estimator, is 0.0233.

```
dat2 <- dat[age <= 7]
```

```
WRS::olsch4(dat2$age, dat2$pep)
```

```
$n
[1] 14

$n.keep
[1] 14

$ci
      Coef. Estimates   ci.lower ci.upper    p-value Std.Error
(Intercept)      0 3.5148814 2.45961431 4.5701484 1.004728e-05 0.48433121
Slope            1 0.2474008 0.03559377 0.4592078 2.570292e-02 0.09721213

$cov
      [,1]      [,2]
[1,] 0.2345767 -0.045083903
[2,] -0.0450839  0.009450198
```

```
WRS::regci(dat2$age, dat2$pep)
```

```
[1] "Duplicate values detected; tshdreg might have more power than tsreg"
[1] "Taking bootstrap samples. Please wait."
```

```
$regci
      ci.low   ci.up Estimate    S.E.    p-value
Intercept 2.55336867 4.7290698 3.0666667 0.6468266 0.01669449
Slope 1    0.02325581 0.4871795 0.3333333 0.1278130 0.03338898
```

```
$n
[1] 14

$n.keep
[1] 14
```

8.16

For the reading data in Table 8.5, verify that the R function `spearci` returns a p-value equal to 0.014 and that `scorci` returns a p-value equal to 0.002. Based on the plot returned by `scorci`, why is it not surprising that these two function give similar results?

```
dat <- data.table::fread(paste0(data.dir, "read_dat.txt"), fill = T, skip = 13)
```

```
spearci( dat[,4], dat[,8])
```

```
$cor.ci
```

```
[1] -0.48330768 -0.01914648
```

```
$p.value
```

```
[1] 0.038
```

```
$cor.est
```

```
[1] -0.2687277
```

```
scorciMC( dat[,4], dat[,8])
```

Attaching package: 'MASS'

The following object is masked from 'package:WRS':


ltsreg

The following object is masked from 'package:formattable':

area

The following object is masked from 'package:dplyr':

select



```
$cor.ci
[1] -0.6900937 -0.1184465
```

```
$p.value
[1] 0.01
```

```
$cor.est
[1] -0.3868361
```

8.17

Given that $b_1 = -1.5$, $n = 10$, $s_{y,x}^2 = 35$, $\sum (X_i - \bar{X})^2 = 140$, assume normality and homoscedasticity and compute a 0.95 confidence interval for slope, β_1 .

```
b1 <- -1.5; n <- 10; se <- sqrt(35 / 140)
alpha <- .05

b1 + qt(c(Lower = alpha/2, Upper = 1 - alpha/2), df = n - 2) * se
```

```
      Lower      Upper
-2.6530021 -0.3469979
```

8.18

Repeat the previous problem, only compute a 0.98 confidence interval.

```
b1 <- -1.5; n <- 10; se <- sqrt(35 / 140)
alpha <- 1 - .98

b1 + qt(c(Lower = alpha/2, Upper = 1 - alpha/2), df = n - 2) * se
```

```
      Lower      Upper
-2.94822972 -0.05177028
```

8.19

Based on results covered in the previous chapters, speculate about why the confidence intervals computed in the last two problems might be inaccurate.

Least squares regression can be negatively impacted by non-normality, heteroscedasticity and outliers.

8.20

Assume normality and homoscedasticity and suppose $n = 30$, $\sum (X_i) = 15$, $\sum (Y_i) = 30$, $\sum (X_i - \bar{X})(Y_i - \bar{Y}) = 30$, $\sum (X_i - \bar{X})^2 = 10$

```
n <- 30; A <- 30; C <- 10;
xbar <- 15 / 30; ybar <- 30 / 30

slope <- A / C
intercept <- ybar - xbar*slope
```

Determine the least squares estimates of the slope and intercept.

$$\beta_1 = 3, \beta_0 = -0.5$$

8.21

Assume normality and homoscedasticity and suppose $n = 38$, $\bar{Y} = 20$, $\sum X_i^2 = 1,922$, $\bar{X} = 7$, $\sum (X_i - \bar{X})(Y_i - \bar{Y}) = 180$, $\sum X_i - \bar{X}^2 = 60$, $s_{X,Y}^2 = 121$.

- Determine the least squares estimate of the slope and intercept.
- Test the hypothesis $H_0 : \beta_0 = 0, \alpha = 0.05$
- Compute a 0.9 confidence interval for β_1

8.22

Assume normality and homoscedasticity and suppose $n = 41$, $\bar{Y} = 10$, $\bar{X} = 12$, $\sum (X_i - \bar{X})(Y_i - \bar{Y}) = 100$, $\sum (X_i - \bar{X})^2 = 400$, $s_{Y.X}^2 = 144$.

- Determine the least squares regression line
- Compute a .9 confidence interval for β_1

8.23

Assume normality and homoscedasticity and suppose $n = 18$, $\beta_1 = 3.1$, $\sum (X_i - \bar{X})^2 = 144$, $s_{X,Y}^2 = 36$.

Compute a 0.95 confidence interval for β_1 .

Would you conclude $\beta_1 > 2$?

8.24

Assume normality and homoscedasticity and suppose $n = 20$, $\beta_0 = 6$, $\sum X_i^2 = 169$, $S_{Y.X}^2 = 25$, $\sum (X_i - \bar{X})^2 = 90$.

Compute a .95 confidence interval for β_0 .

8.25

Given the following quantities, find the sample correlation coefficient, r , and test $H_0 : \rho = 0$ at the indicated value for α .

a.) $n = 27$, $\sum (Y_i - \bar{Y})^2 = 100$, $\sum (X_i - \bar{X})^2 = 625$, $\sum (X_i - \bar{X})(Y_i - \bar{Y})^2 = 200$, $\alpha = 0.01$

b.) $n = 5$, $\sum (Y_i - \bar{Y})^2 = 16$, $\sum (X_i - \bar{X})^2 = 25$, $\sum (X_i - \bar{X})(Y_i - \bar{Y})^2 = 10$, $\alpha = 0.05$

8.26**8.27****8.28****8.29****8.30****8.31****8.32****8.33****8.34****8.35****8.36**